

# Week 1 Notes - ML

Start Date @27 October 2025

These are notes that I've taken whilst reading the notes from Chapter 1 and 2 of the Machine Learning with PyTorch & Scikit-Learn book by Sebastian Raschka and others. The notes are roughly the same as the hand-written notes, however sometimes there are minor changes.

## Types of Machine Learning

- **Supervised Learning** → Labelled data, direct feedback, predict outcome/future
- **Unsupervised Learning** → Unlabelled data, no feedback, find hidden structure in data.
- **Reinforcement Learning** → Decision process, reward system, learn series of actions

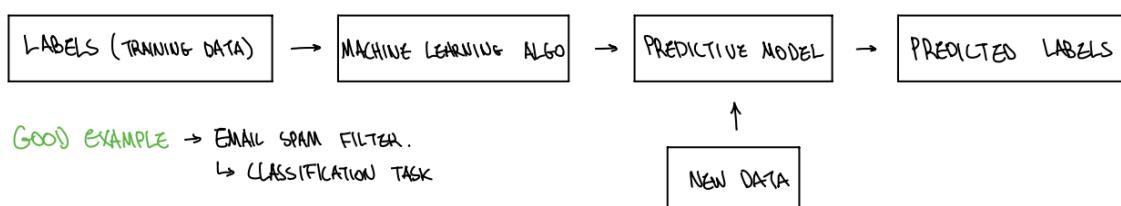
## Supervised Learning



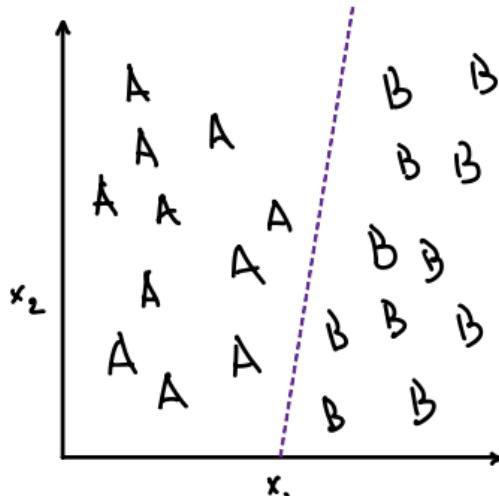
Model learns from labelled training data → Makes future predictions → aka.  
"Labelled Learning"

Good example → Email spam filter (Classification task)

## Classification



- Sub-category of supervised learning
- **Aim** → Predict categorial class labels of new instances, based on past examples



Classification (Class A or Class B)

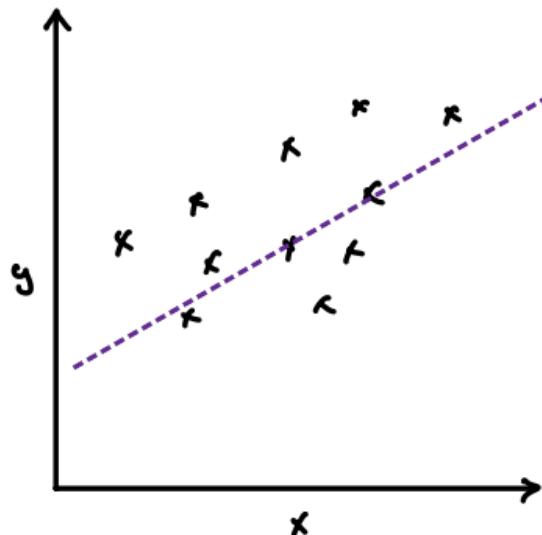
- Email spam filter → **Binary Classification** task, is the email a spam or not a spam email
- Supervised learning trained model → assign any class label to new data
  - **Multiclass Classification**
    - Handwritten character recognition → "A", "B", "C"
    - However, this model wouldn't recognise numbers for example, because it was only trained on letters

## Regression



Variables are called **features, target variables**

- Finding the relationship between variables to predict an outcome.

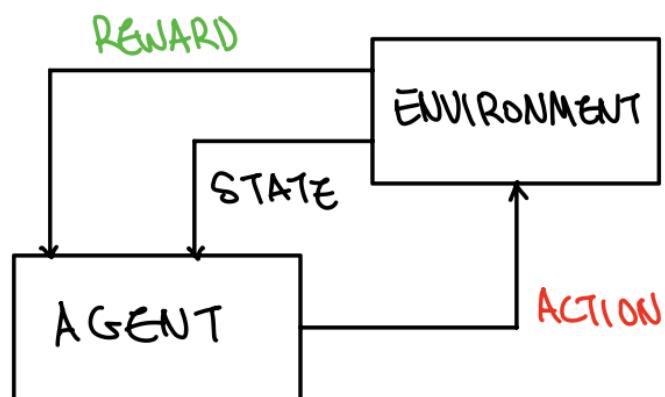


Regression model

- E.g. → Finding correlation between time spent studying vs. exam results (training dataset) → predict future scores of students depending on the time they spent studying

## Reinforcement Learning

- Develop agent → Improves performance based on interaction with environment
- Trial-and-error → machine learns with **rewards** system → chooses a series of actions to maximise rewards
  - Feedback is immediate or delayed feedback
- E.g. → Chess program → Win or lose with different moves → understands whilst playing the player, improves system as it goes on.



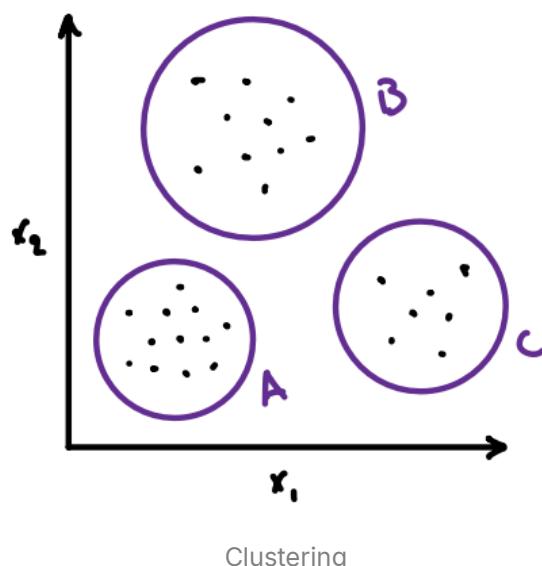
Reinforcement Learning

# Unsupervised Learning

- Model trains with unlabelled data
- Explore structure of data to extract data without guidance from known outcome variables

## Clustering

- Organise information into subgroups (clusters) → without prior information about group
- Sometimes called **Unsupervised Classification**
- Good technique for structuring information + finding relationships from data



## Dimensional Reduction

- Often we work with high dimensionality (larger number of features)
- Compresses data onto smaller dimensional subspace while retaining most of the relevant information
- **Commonly used** → feature reprocessing to remove noise from data
- Can be useful for data visualisation (sometimes)

## Terminology

- **Training example** → Row in a table representing dataset used to teach a model how to make predictions

- **Training** → Model fitting
- **Feature (x)** → Input/variable
- **Target (y)** → Output/outcome
- **Loss Function**
  - Cost function/error function
  - “**Loss**” → Loss measured for single data point
  - “**Cost**” → Compute loss over the entire dataset (average or summed)

# Machine Learning Systems

## Preprocessing



**One of the most crucial steps** → shape data for learning algorithms

- Select features → should be the same scale
- Dimensional Reduction → Highly correlated features, therefore redundant to certain degree
  - **Advantage** → Less storage space required + algorithm can run faster
- Divide dataset to test new data against trained data



80:20 → 80 test, 20 validation (testing)

## Training, Selecting and Evaluating Models

- One model isn't necessarily good for all applications
- **Essential** to train multiple models and train them to evaluate which works best
- Common metric to measure performance → **Classification accuracy**
- “How do we know which model performs well on final data?” → **Cross-validation**
- Cross-validation → Further divide dataset into training + validation subset
- Frequent use → Hyperparameter fine-tuning
- **Generalization Error** → How accurately will the training model perform with test dataset (predict outcome)

- Same procedures applied to training dataset applied to test dataset → otherwise model will **overfit**



**Overfitting** → Machine learning model learns from the training data too well (includes noise), resulting in poor performance on new, unseen data.

## Classification models

- **Perceptron**
- **Adaptive Linear Neurons (Adaline)**

## Artificial Neurons

- First computational model of a neuron → **McCulloh-Pitts** (1943)
  - Concept made by Warren McCulloh + Walter Pitts
- First concept of perceptron learning rule based on MCP → Frank Rosenblatt (1957)
  - Proposed an algorithm that learns optimal weight → multiplied with input features → order to make decision if neuron fires or doesn't
  - **Current Day** → In Regression + Classification
    - This algorithm could be used to predict whether a new data point belongs to one class or other
- **Binary Classification** → whether data is 0 or 1

$$w = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix} \quad (1) \quad x = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix} \quad (2) \quad \sigma(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{otherwise} \end{cases}$$

## Linear Algebra

- Abbreviate the sum of products of value  $\textcolor{red}{x}$  and  $\textcolor{red}{w}$  → vector dot product

$$a = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix}, \quad b = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \quad (3) \quad \text{Transpose } a \text{ to } a^T \quad a^T = [a_1 \quad a_2 \quad a_3]$$

$$\mathbf{a}^T \mathbf{b} = \sum_{i=1} a_i b_i = a_1 \cdot b_1 + a_2 \cdot b_2 + a_3 \cdot b_3$$

- Example:

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}^T = \begin{bmatrix} 1 & 3 & 5 \\ 2 & 4 & 6 \end{bmatrix}$$

## Perceptron Learning Rule

- Reductionist approach to mimic how single neuron works in the brain.
- Algorithm summarised:
  1. Init weight + bias  $\rightarrow$  until 0 or small number is reached
  2. For each training example  $x^{(i)}$ :
    - a. Compute output  $\hat{y}^{(i)}$
    - b. Update weight + bias unit
- Output is class label predicted by unit step function

$$w_j := w_j + \Delta w_j \quad b := b + \Delta b$$

$$\Delta w_j = \eta \left( y^{(i)} - \hat{y}^{(i)} \right) x_j^{(i)} \quad \Delta b = \eta \left( y^{(i)} - \hat{y}^{(i)} \right)$$

$w_j = x_j$  in dataset

$\eta$  = Learning rate (Typically constant between 0.0 and 1.0)

$y^{(i)}$  = True class label of  $i$ th training example

$\hat{y}^{(i)}$  = Predict class label

- Bias unit + weight  $\rightarrow$  constantly updated simultaneously

$$\Delta w_1 = \eta \left( y^{(i)} - \text{Output}^{(i)} \right) x_1^{(i)}$$

$$\Delta w_2 = \eta \left( y^{(i)} - \text{Output}^{(i)} \right) x_2^{(i)}$$

$$\Delta b = \eta \left( y^{(i)} - \text{Output}^{(i)} \right)$$

- **Perceptron predicts correctly** → Bias unit + weight remain unchanged → value are 0:

$$y^{(i)} = 0, \quad \hat{y}^{(i)} = 0, \quad \Delta w_j = \eta(0 - 0)x_j^{(i)} = 0, \quad \Delta b = \eta(0 - 0) = 0$$

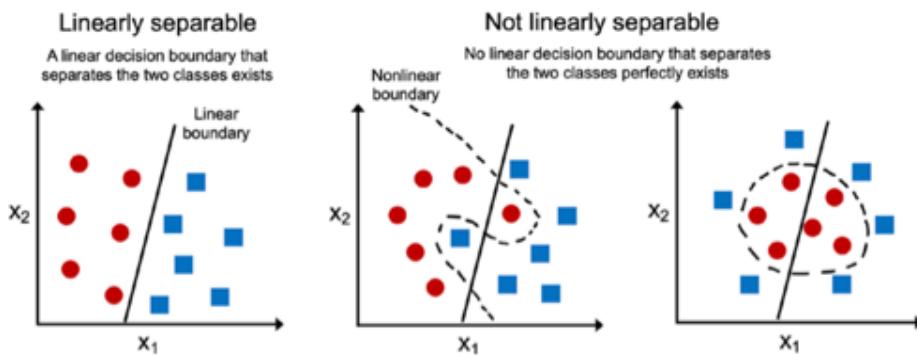
$$y^{(i)} = 1, \quad \hat{y}^{(i)} = 0, \quad \Delta w_j = \eta(1 - 1)x_j^{(i)} = 0, \quad \Delta b = \eta(1 - 1) = 0$$

- **Wrong prediction** → Leans positive or negative target class:

$$y^{(i)} = 0, \quad \hat{y}^{(i)} = 0, \quad \Delta w_j = \eta(1 - 0)x_j^{(i)} = \eta x_j^{(i)}, \quad \Delta b = \eta(0 - 0) = \eta$$

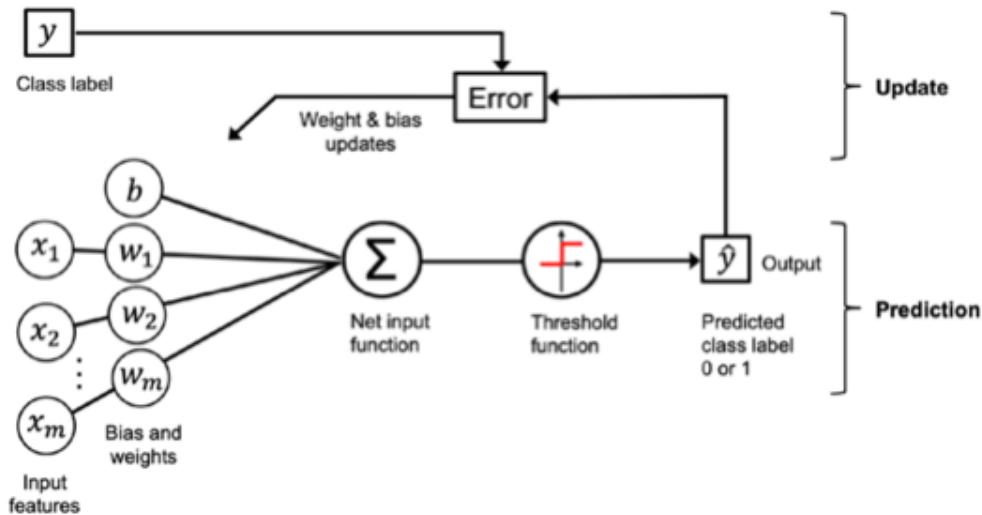
$$y^{(i)} = 1, \quad \hat{y}^{(i)} = 0, \quad \Delta w_j = \eta(1 - 1)x_j^{(i)} = -\eta x_j^{(i)}, \quad \Delta b = \eta(1 - 1) = -\eta$$

- Convergence of perceptron only guaranteed → two classes are linearly separable



- If two classes can't be separated by linear boundary → set maximum number of passes over training dataset (epoch)

## General Concept of Perceptron



- **Input examples ( $x$ ) + bias ( $b$ ) + weights ( $w$ )**  $\Rightarrow$  compute net input
  - Helps shift decision boundary
- **Threshold Function**  $\rightarrow$  Checks if sum  $x$  crosses certain threshold

$$\text{Decides Output} = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{otherwise} \end{cases}$$

- **Learning Phase**
  - When prediction is wrong  $\rightarrow$  adjusts weights to get closer to correct output
  - Continues until perceptron classifies all training examples correctly

## Perceptron on Iris Dataset

- Only use two flower classes  $\rightarrow$  Perceptron is a binary classifier
- Perceptron algorithm can be extended  $\rightarrow$  One-versus-All (OvA) technique
- **One-versus-All (OvA)**  $\rightarrow$  Multi-class Classifier
  - Sometimes called One-versus-Rest (OvR)
  - Extend Binary Classifier  $\rightarrow$  multi-class problem
  - Train one **classifier per class**  $\rightarrow$  treated as positive + examples from all other classes are negative
  - New unlabelled data  $\rightarrow$  use  $n$  classifiers,  $n =$  Number of class labels  $\rightarrow$  assign to highest confidence to particular instance

- In case of perceptron → OvA → Choose class label that is associated with largest absolute net input value

## Perceptron Convergence



Convergence → Biggest problem for the perceptron

- If classes cannot be separated → weights will never stop updating unless we set maximum number of epochs

## Adaptive Linear Neurons (Adaline)

- Single-layer Neural Network → Adaptive Linear Neurons (Adaline) 1960

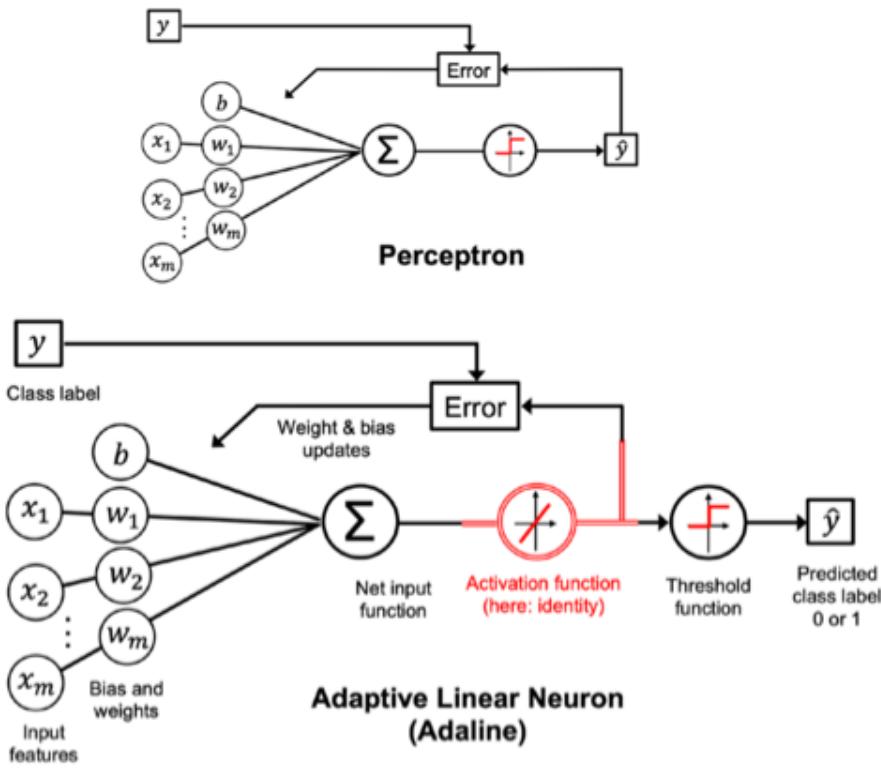
- Bernard Widrow + Ted Hoff



Key concepts of defining and minimising continuous loss function

- Lay the ground works for understanding other machine learning algorithms for classification. Examples:
  - Logistic Regression
  - Support Vector Machines
  - Multi-layer Neural Networks
  - Also → Linear Regression Model
- Key difference with Rosenblatt's perceptron → Weights updated based on linear activation function vs. unit step function in Rosenblatt's perceptron
- Linear activation function =  $\sigma(z)$  → Identity function of net input  $\sigma(z) = z$

## Main Concept of Adaline



- **Input examples (xx) + bias (bb) + weights (ww) →→ compute net input**
  - Sums up weighted inputs and bias for every training example
- **Linear Output →→ Direct continuous value (no threshold yet)**
  - Output is not just “yes/no,” but any real number
- **Learning Phase**
  - Calculates error: how far the output is from the correct value
  - Uses gradient descent/delta rule to adjust weights and bias to minimise this error
- **Final Decision (Threshold)**
  - After training, applies a threshold function to convert output to class:

$$\text{Decides Output} = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{otherwise} \end{cases}$$

## Minimising Loss Function with Gradient Descent

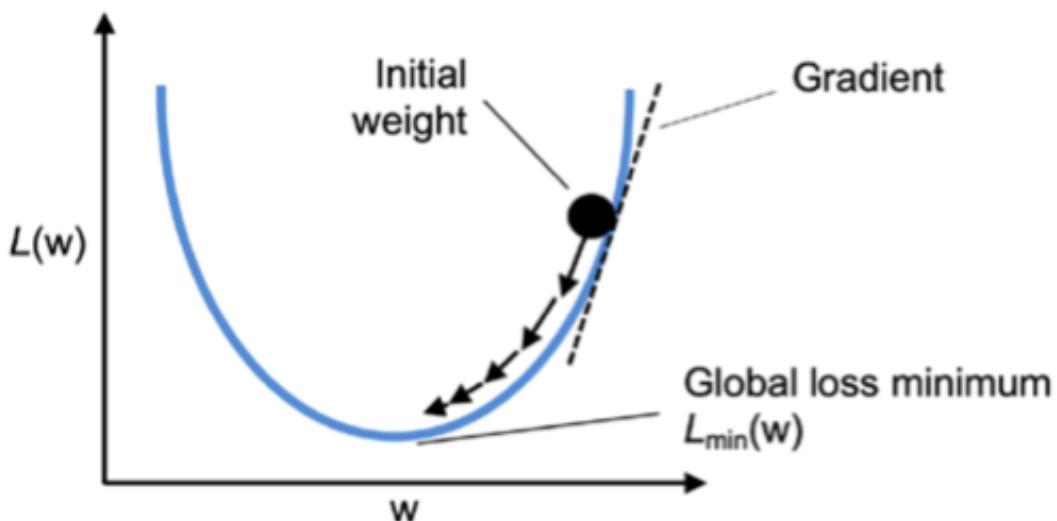
- Key ingredient for supervised learning → **Objective Function**

- **Objective Function** → Loss or cost function that we want to minimise
- Adaline → Loss function ( $L$ ) → Learn model parameters as Mean Square Error (MSE) → between calculated + true class label

$$L(w, b) = \frac{1}{2n} \sum_{i=1}^n \left( y^{(i)} - \sigma(z^{(i)}) \right)^2$$

- $\frac{1}{2}$  for our own convenience → easier to derive the gradient of loss function
- Advantage:
  - Loss function becomes differentiable
  - It's convex → Can use Gradient Descent → Find weights that maximise loss function to classify examples in a dataset

## Gradient Descent



- "Climbing down hill" → Until loss or global loss minimum is reached
- Each iteration → step opposite direction of gradient
  - Step determined by learning rate + slope of gradient
- Opposite direction of  $\nabla L(w, b)$  :  $w : w + \Delta w, b : b + \Delta b$
- $\Delta w, \Delta b$  defined as negative gradient multiplied by learning rate  $\eta$

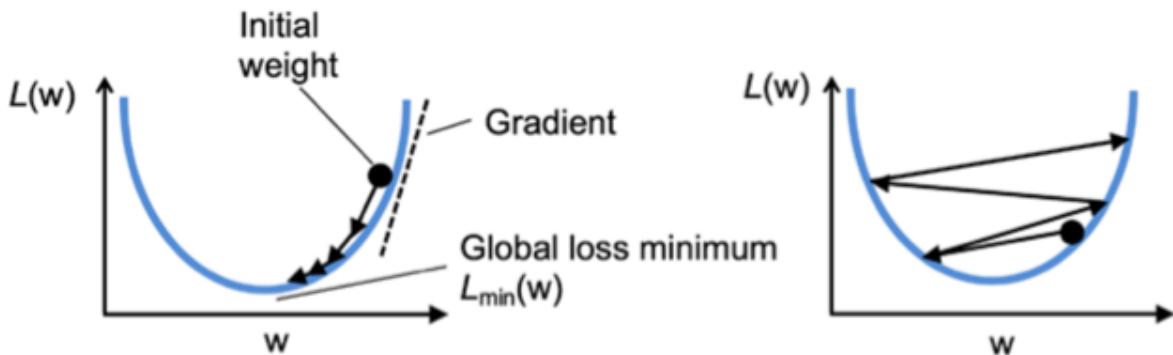
$$\Delta w = -\eta \nabla_w L(w, b), \quad \Delta b = -\eta \nabla_b L(w, b)$$

- Gradient of loss function → partial derivative of loss function → respect to each weight,  $w_j$

$$\frac{\partial L}{\partial w_j} = -\frac{2}{n} \sum_i (y^{(i)} - \sigma(z^{(i)})) x_j^{(i)}$$

- Partial derivative of loss respected to bias

$$\begin{aligned}\frac{\partial L}{\partial w_j} &= \frac{\partial}{\partial w_j} \frac{1}{n} \sum_i (y^{(i)} - \sigma(z^{(i)}))^2 = \frac{1}{n} \frac{\partial}{\partial w_j} \sum_i (y^{(i)} - \sigma(z^{(i)}))^2 \\ &= \frac{2}{n} \sum_i (y^{(i)} - \sigma(z^{(i)})) \frac{\partial}{\partial w_j} (y^{(i)} - \sigma(z^{(i)})) \\ &= \frac{2}{n} \sum_i (y^{(i)} - \sigma(z^{(i)})) \frac{\partial}{\partial w_j} \left( y^{(i)} - \sum_j (w_j x_j^{(i)}) + b \right) \\ &= \frac{2}{n} \sum_i (y^{(i)} - \sigma(z^{(i)})) \frac{\partial}{\partial w_j} \left( y^{(i)} - \sum_j (w_j x_j^{(i)}) + b \right) \\ &= \frac{2}{n} \sum_i (y^{(i)} - \sigma(z^{(i)})) (-x_j^{(i)}) = -\frac{2}{n} \sum_i (y^{(i)} - \sigma(z^{(i)})) x_j^{(i)}\end{aligned}$$



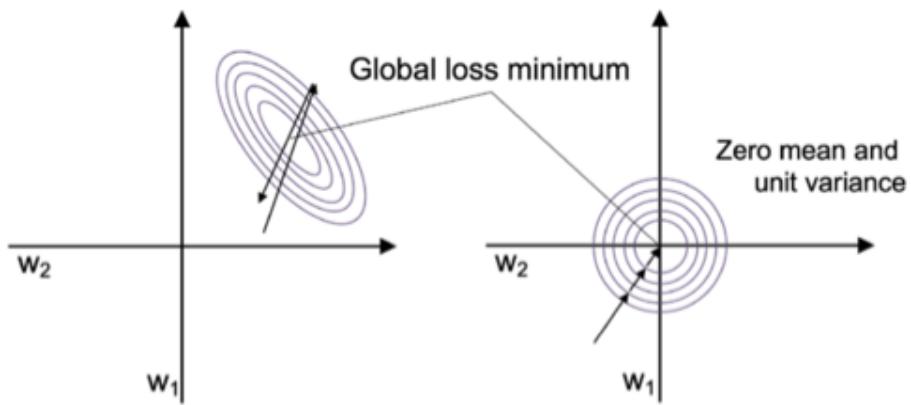
- Left Image → well-chosen learning rate (hyperparameters) where loss decreases gradually
- Right Image → Overshot global minimum → Learning rate is too large

## Standardisation

- Normalisation that helps gradient descent learning to converge more quickly.  
**Dataset not normally distributed**
  - Shifts the mean of each feature → centred at zero + each features standard deviation = 1

$$x'_j = \frac{x_j - \mu_j}{\sigma_j}$$

- Helps gradient descent learning → Easier to find learning rate that works well for weights + bias
- Great for features with different scales → stabilises training → optimiser can go through fewer steps for optimal solution



- Left unscaled features
- Right standardised features

## Stochastic Gradient Descent (SGD)

- Aka. Iterative or online gradient descent
- Parameters are updated incrementally for each training example over sum of accumulated errors over all training examples

$$\Delta w_j = \eta(y^{(i)} - \sigma(z^{(i)}))x_j^{(i)}, \quad \Delta b = \eta(y^{(i)} - \sigma(z^{(i)}))$$

- SGD considered an approximation to gradient descent → typically reaches convergence faster → more frequent weight updates
- Error surface is noisier than gradient descent → advantage for SGD → escape shallow minima more readily → nonlinear loss function
- SGD → present training data in random order → shuffle data every epoch → prevents cycles
- SGD - adaptive learning rate → decreases over time

$$\frac{c_1}{\text{number of iterations} + c_2}$$

- Advantage of SGD → Can do Online Learning
  - Model trained with new data on the fly
  - Useful for large amounts of data
  - System can immediately adapt to changes

## Mini-batch Gradient Descent

- Full batch gradient descent to smaller subsets of training data
- Advantage over full batch → Convergence reached faster (finds optimal weights and biases faster)