

# Week 2 - ML

Start Date @10 November 2025

★ Chapter 3

## Choosing Classification Algorithm

- Takes practice and time

| "No free lunch theorem" - David H. Wolpert

- There isn't one classifier that works best for all scenarios
- Choose a handful of models, then select one that works → model training very dependant on data over training model
- **5 steps fro supervised training:**
  1. Select features + collect labelled training data
  2. Choose performance metric
  3. Choose learning algorithm + training model
  4. Evaluate performance of model
  5. Fine-Tune model
- Feature selection, pre-processing, etc.. is important too but covered later in the book

## Scikit-Learn - Perceptron

- User-friendly + optimised implementation of several classification algorithms
- `train_test_split` → split data into test + training data
- Many machine learning + optimisation algorithms → require feature scaling for optimal performance (e.g. Gradient Descent)

## Logistic Regression + Conditional Probabilities



Logistic Regression → Classification model → performs well on linearly separable classes

- **One of the most widely used algorithms for classification** → easy to implement and scale
- Linear model for binary classification
- Logistic Regression for multiple classes → **Multinomial Logistic Regression or Softmax regression**

## Main mechanics

---

### Odds

Odds in favour of particular event.

$$\frac{p}{(1 - p)}$$

p = probability of positive event

- positive event = event we want to predict. Doesn't necessarily mean that it's positive

### Logit Function

$$\text{logit}(p) = \log \frac{p}{(1 - p)}$$


- $\log$  → Natural log
- Takes input value in the range of 0 to 1

$$\text{logit}(p) = w_1x_1 + \dots w_mx_m + b = \sum_{i=j} w_jx_j + b = w^T x + b$$

- Logit function maps probability to real-number range → consider inverse of function to map back to [0,1] range

### Inverse of Logit



- Logistic Sigmoid Function → abb. Sigmoid Function → S shape



Logistic Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- $z = \text{net input} \rightarrow z = w^T x + b$



Logistic Regression compared to Adaline

- $\sigma(z) \rightarrow 1$  if  $z \rightarrow \infty$  because  $e^{-2}$  becomes very small for large values
- $\sigma(z) \rightarrow 0$  if  $z \rightarrow -\infty$  as a result of an increasingly large denominator
- Conclusion is that Sigmoid takes real-number values as input and transforms them into values in range 0 to 1. Intercept  $\sigma(0) = 0.5$



Logistic Regression used in weather forecasting and also popular in medicine

## Model Weights Via Logistic Loss Function

- Minimise MSE to learn parameters for Adaline classification model

$$L(w, b | x = p(y|x; w, b)) = \prod_{i=1}^n p(y^{(i)}; w, b) = \prod_{i=1}^n (\sigma(z^{(i)}))^{y^{(i)}} (1 - \sigma(z^{(i)}))^{1-y^{(i)}}$$

**In practice → we would use log-likelihood**

$$L(w, b | x) = \text{Log}L(y|x; w, b) = \sum_{i=1}^n [y^{(i)} \log(\sigma(z^{(i)})) + (1 + y^{(i)}) \log(1 - \sigma(z^{(i)}))]$$

### ▼ Logistic Loss Function

$$L(\theta) = \prod_{i=1}^n P(y^{(i)}|x^{(i)}; \theta)$$

Log-Likelihood

$$\text{log}L(\theta) = \sum_{i=1}^n \log P(y^{(i)}|x^{(i)}; \theta)$$

1. Applying Log Function reduces potential numeric underflow
  - Working with probability in range 0-1
  - Sometimes probabilities can get really close to 0 or 1.
  - $\log(p)$  makes very tiny probabilities manageable + numerically stable
    - Improves reliability for model training process
2. Convert product of factors into summation of factors → Easier to compute derivatives

- e.g.  $\log(a \times b \times c) = \log(a) + \log(b) + \log(c)$
- Taking log of products allows you to write the loss as a sum, making gradient calculation for weights simpler + stable



### Foundation of Logistic Regression → Likelihood Function

- Goal → Maximise the likelihood → Done with Gradient Ascent
  - **However**, can use gradient descent with negative log-likelihood




Machine Learning frameworks are designed for minimising → it's common to switch to negative and would be the better option too.

$$L(w, b) = \sum_{i=1}^n [-y^{(i)} \log(\sigma(z^{(i)})) + (1 + y^{(i)}) \log(1 - \sigma(z^{(i)}))]$$

$$L(\sigma(z), y; w, b) = -y \log(\sigma(z)) - (1 - y) \log(1 - \sigma(z))$$

$$L(\sigma(z), y; w, b) = \begin{cases} -\log(\sigma(z)) & \text{if } y = 1 \\ -\log(1 - \sigma(z)) & \text{if } y = 0 \end{cases}$$



A plot of the loss function used in logistic regression

**Main Point → We penalise wrong predictions with increasingly larger loss**

## Converting Adaline Implementation Into Algorithm For Logistic Regression

- Compute loss of classifying per epoch
- Swap linear activation with Sigmoid
  - Switch the first two and it becomes Logistic Regression
- Fit Logistic Regression model → Only works for binary classification tasks

## Gradient Descent Learning Algorithm For Logistic Regression

- Adaline, weights + bias didn't change for Logistic Regression → Gradient Descent similar for Logistic Regression
- Partial derivative of log-likelihood with respect to  $j$ th term:

$$\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial w_j} \quad \text{WHERE } a(z) = \frac{1}{1+e^{-z}}$$

$$\frac{\partial L}{\partial a} = \frac{a-y}{a \cdot (1-a)} \quad \left[ \begin{array}{l} \frac{\partial L}{\partial z} = a - y \\ \frac{\partial z}{\partial w_j} = x_j \end{array} \right] \quad \left[ \begin{array}{l} \frac{\partial L}{\partial w_j} = (a-y)x_j \\ \frac{\partial L}{\partial w_j} = -(y-a)x_j \end{array} \right]$$

→ TAKE STEP OPPOSITE OF GRADIENT → FLIP  $\frac{\partial L}{\partial w_j} = -(y-a)x_j + \text{LEARN RATE } \eta$


$$w_j := w_j + \eta(y-a)x_j$$


$$b := b + \eta(y-a)$$

**WEIGHT + BIAS = ADALINE WEIGHT + BIAS**

## Logistic Regression with Scikit-Learn

- Supports multi-class Logistic Regression off the shelf






Multi-class = Multinomial (Default for Scikit-learn Logistic Regression)

Convex Optimisation → minimising convex loss functions (logistic regression) →  
**recommended to use more advanced algorithm over SGD**

## Overfitting via Regularisation

- Common problem for machine learning → Doesn't generalise well with unseen data
- Overfitting → High variance model → **caused by too many parameters (too complex)**
- Underfitting → High bias → **model not complex enough**



Examples of underfitted, well-fitted, and overfitted models

## Bias-Variance Trade-Off

### What is Bias?

- Measures how far off the predictions are from the correct values
- High-bias model makes strong assumptions about data but underfits - **can't figure out the underlying pattern well**

### What is Variance

- Measures consistency of model prediction
- High-variance  $\rightarrow$  model is sensitive to specific training data (fits noise + small data fluctuation)

### The Trade-Off

- Gotta find the sweet spot between the two
- One way to find good bias-variance trade-off  $\rightarrow$  **tune complexity of model via regularisation**



Regularization  $\rightarrow$  Very useful for handling collinearity, filtering out noise from data + eventually preventing overfitting

- Add additional information to penalise extreme parameter values  $\rightarrow$  most common  $\rightarrow$  **L2-Regularization**
- Regularisation  $\rightarrow$  feature scaling and standardisation  $\rightarrow$  are all important
  - For this to work, features must be on comparable scales

$$\frac{\lambda}{2n} = \frac{\lambda}{2n} \sum_m^{j=1} w_j^2$$

- $\lambda$  = Regularisation parameter
- 2 = Scaling factor → cancels computing loss function
- $n$  = Sample size added to scale regularisation term similar to loss

### **Loss function for logistic regression + regularisation:**

$$L(w, b) = \sum_{i=1}^n [-y^{(i)} \log(\sigma(z^{(i)})) + (1 + y^{(i)}) \log(1 - \sigma(z^{(i)}))] + \frac{\lambda}{2n} \|w\|^2$$

### **Partial derivative of unregularised loss:**


$$\frac{\partial L(w, b)}{\partial w_j} = \left( \frac{1}{n} \sum_{i=1}^n (\sigma(w^T x^{(i)}) - y^{(i)}) x_j^{(i)} \right)$$

### **Adding regularisation to derivative of loss:**

$$\frac{\partial L(w, b)}{\partial w_j} = \left( \frac{1}{n} \sum_{i=1}^n (\sigma(w^T x^{(i)}) - y^{(i)}) x_j^{(i)} \right) + \frac{\lambda}{n} w_j$$



We can control how closely we fit the data with  $\lambda$  → If we increase  $\lambda$ , increase regularisation strength.



The impact of the inverse regularization strength parameter  $C$  on L2 regularized model results




### Why do we not strongly regularie all models if it reduces overfitting?

If regularisation is too high → weights coefficient approaches 0 and therefore our model will underfit. (look at fig above)

## Maximum Margin Classification With Support Vector Machines (SVM)

- Powerful + widely used algorithm → **Support Vector Machine (SVM)** → **considered extension of perceptron**
- SVM → maximise margin → Distance between the **separating hyperplane (decision boundary)** and **training examples closest to hyperplane (support vectors)**



SVM maximizes the margin between the decision boundary and training data points

## Maximum Margin Intuition


- Larger margins  $\rightarrow$  tend to have a lower generalization error



SVM looks simple  $\rightarrow$  math is complex  $\rightarrow$  need to know **constrained optimisation**


## Slack Variable

- Introduced by Vladimir Vapnik in 1995  $\rightarrow$  this led to **soft-margin classification**
- Strict constraints  $\rightarrow$  this force the SVM to find a hard boundary. This fails if data isn't cleanly separated
- Relaxing constraints  $\rightarrow$  **this would be adding slack variables**  $\rightarrow$  some points can be misclassified, but will get penalised for each mistake
- SVM optimisation seeks balance  $\rightarrow$  wide margin (good separation) + few misclassification (low slack penalty)
- To use slack variables**  $\rightarrow$  introduce variable  $C$   $\rightarrow$  hyperparameter for controlling penalty for misclassification



The impact of large and small values of the inverse regularization strength  $C$  on classification

- Large  $C \rightarrow$  the larger the error penalties
- Smaller  $C \rightarrow$  less strict about misclassification error
- $C$  controls width of margin, therefore toning bias-variance tradeoff



SVM's decision regions

## Logistic Regression vs. SVM

- Logistic Regression and SVM in practical classification yield similar results
- Logistic Regression → Maximises the conditional likelihood
- SVM → Cares about points closest to decision boundary (support vectors)



Logistic Regression is the simpler model, it's easier to implement and is easily updatable with newer data.