

# Week 3 - ML

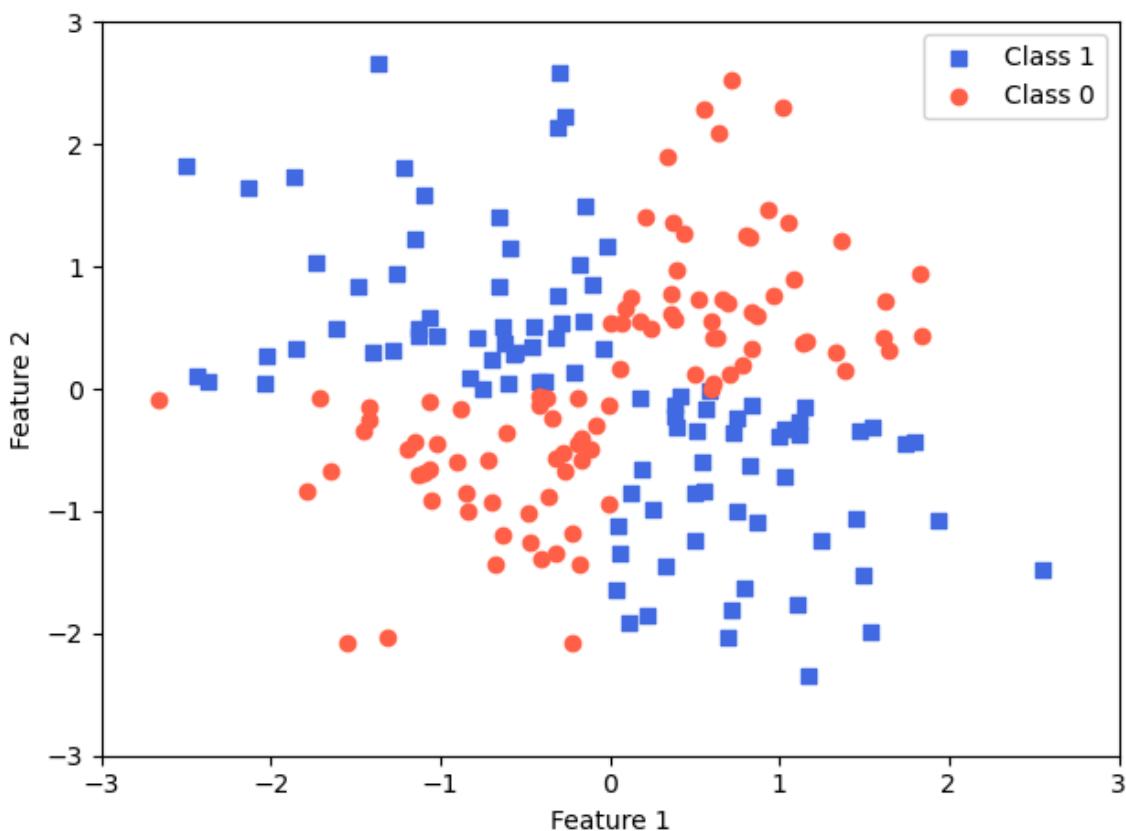
Start Date @17 November 2025

Weeks Week 3

⭐ still chapter 3

## Kernel Methods for Linearly Inseparable Data

- SVMs can be kernelized for non-linear classification
- Kernel SVM → Most common variant of SVM



- XOR dataset with noise

- Can't separate examples from negative + positive classes
- Idea of Kernel Method → Create non-linear combinations of original features to project them onto higher-dimensional space via mapping function ( $\phi$ )

$$\phi(x_1, x_2) = (z_1, z_2, z_3) = (x_1 \cdot x_2, x_1^2 + x_2^2)$$

## Using Kernel Trick to Find Separating Hyperplanes In A High-dimensional Space

- Non-linear problem → Mapping Function ( $\phi$ ) → Train Linear SVM to classify data in new feature space →  $\phi$  (Mapping Function) → Transform new unseen data to classify it using Linear SVM.
- However, Mapping Function is computationally very expensive, especially high-dimensional data
  - Use Kernel Trick
- Quadratic programming to train SVM → replace dot product  $x^{(i)}x^{(j)} = \phi(x^{(i)})^T\phi(x^{(j)})$
- To save expensive dot product → we define kernel function

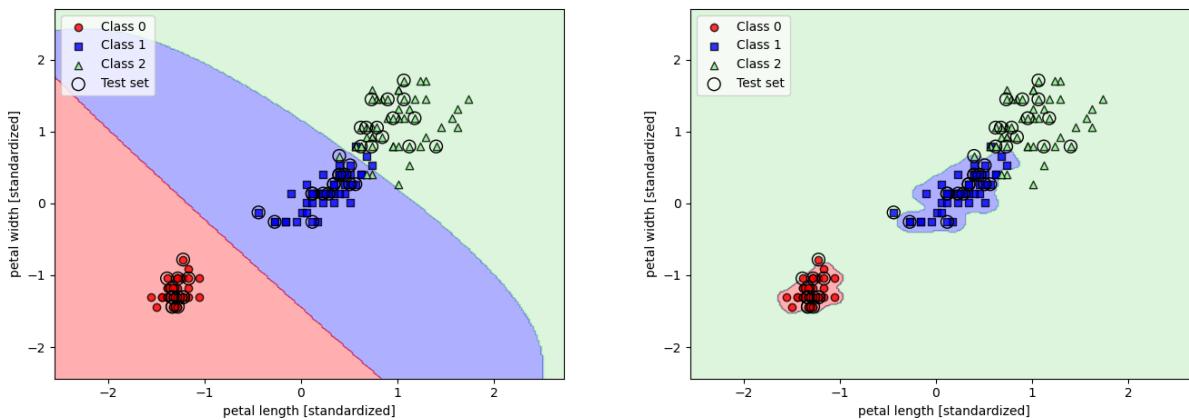
$$K(x^{(i)}, x^{(j)}) = \phi(x^{(i)})^T\phi(x^{(j)})$$

- One of the most widely used Kernels → Radial Basis Function (RBF) → Called Gaussian Kernel

$$K(x^{(i)}, x^{(j)}) = \exp\left(-\frac{\|x^{(i)} - x^{(j)}\|^2}{2\sigma^2}\right)$$

Simplified →  $K(x^{(i)}, x^{(j)}) = \exp(-\gamma\|x^{(i)} - x^{(j)}\|^2)$

- $\gamma = \frac{1}{2\sigma^2}$  → Free parameter to be optimised
- Kernel → Interpreted as similarity function between pair of examples
- Minus sign → Inverts distance into similarity score → falls between range (0,1)



$\gamma$  is small  $\rightarrow$  RBF Kernel SVM model relatively soft

$\gamma$  is high  $\rightarrow$  Fits data but will have high generalization error



$\gamma$  plays an important role in controlling overfitting or variance when algorithm is too sensitive to fluctuation in training dataset

## Decision Tree Learning

- Decision Tree Classifier  $\rightarrow$  Attractive model when we care about **interpretability**
  - Break down data by making decisions based on questions  $\rightarrow$  shape like a tree
- Model learns class labels from features by series of questions asked
- Start at root  $\rightarrow$  split data on feature that results in largest **Information Gain (IG)**
- Training examples in each node belong to the same class
  - Splitting happens until all leaves are “pure”
- We want to **prune** tree by setting limits  $\rightarrow$  tree can become very deep  $\rightarrow$  **prone to overfitting**

## Maximising Information (IG)

- Splitting tree at most informative features  $\rightarrow$  define objective function  $\rightarrow$  optimise algorithm

$$IG(D_p, f) = I(D_p) - \frac{N_{left}}{N_p} I(D_{left}) - \frac{N_{right}}{N_p} I(D_{right})$$

- Three impurity measures or splitting criteria used in binary decision trees:
  - Gini Impurity ( $I_G$ )
  - Entropy ( $I_n$ )
  - Classification Error ( $I_E$ )

## Entropy

- For all non-empty classes ( $P(i|t) \neq 0$ )

$$I_H(t) = - \sum_{i=1}^c P(i|t) \log_2 P(i|t)$$

- $P(i|t)$  = proportion of examples that belong to  $i$  for particular node,  $t$
- Entropy = 0 if all examples at a node belong to same class
- Entropy = maximal - uniform class distribution

## Gini Impurity

- Criterion to minimise the probability of misclassification:

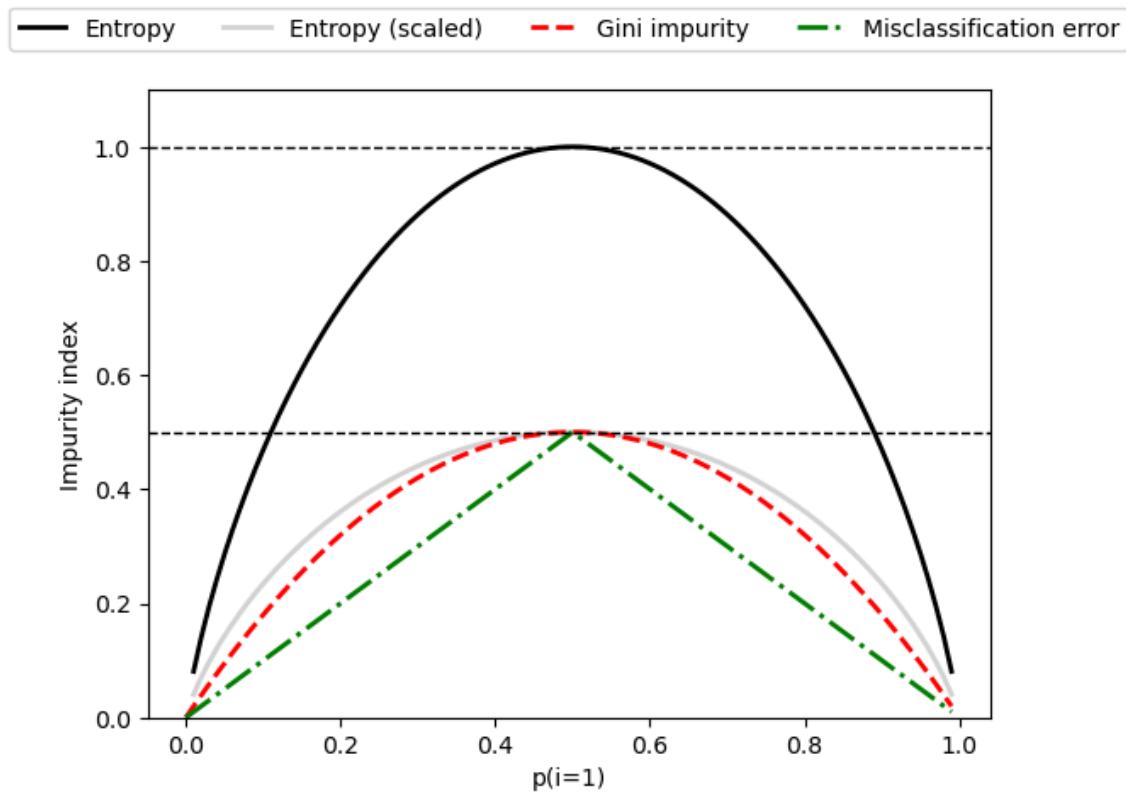
$$I_G(t) = \sum_{i=1}^c P(i|t)(1 - P(i|t)) = 1 - \sum_{i=1}^c P(i|t)^2$$

- Gini Impurity = maximal if all the classes are perfectly mixed
- In practice both Gini Impurity and Entropy → **Typically yield very similar results**
- Not worth evaluating trees using different impurity criteria → Experiment rather with different pruning cut-offs

## Classification Error

$$I_E(t) = 1 - \max\{P(i|t)\}$$

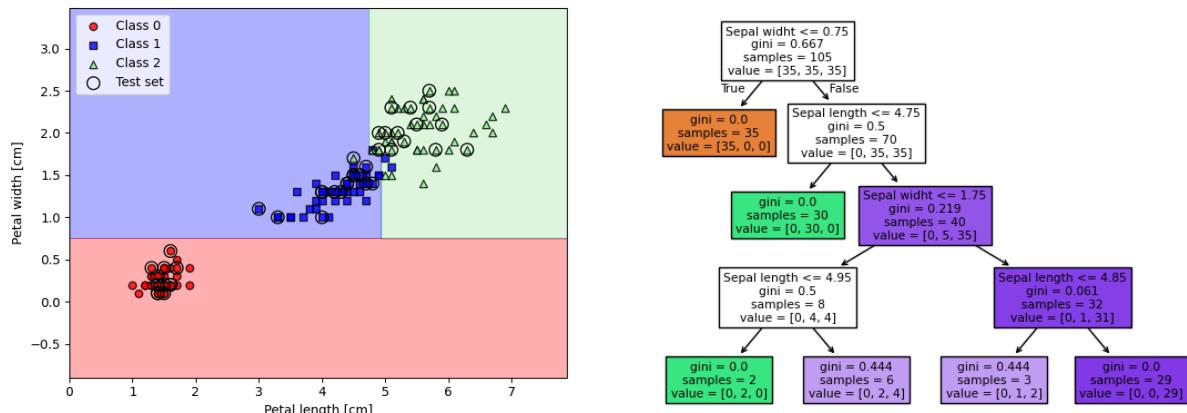
- Useful for pruning → not recommended for growing decision trees → less sensitive to changes in class probabilities of the node



- Different impurity indices for different class-membership probabilities in range 0,1

## Building a Decision Tree (Scikit-Learn)

- Decision Trees can build complex decision boundaries → dividing feature space into rectangles



- Scikit-Learn Decision Tree → max depth of 4, Gini Impurity
- Feature scaling not required for decision trees
- Can visualise decision tree model after training with Scikit-Learn

## Combining Multiple Decision Trees via Random Forest

- Ensemble method → became popular → **Good classification performance + Robustness towards overfitting**
- Random Forest → **Good Scalability + Easy to use**
  - Ensemble of Decision Trees
  - Idea behind Random Forest



Average multiple (deep) Decision Trees that individually suffer from high variance  
 Build a more robust model → better generalisation performance + less susceptible for overfitting

### Four Simple Steps:

1. Draw random **bootstrap** sample size  $n$  (randomly chose a  $n$  examples from training dataset with replacement)
2. Grow Decision Tree for bootstrap sample
  - Each node:
    - a. Randomly select  $d$  features without replacement
    - b. Split node using Objective Function
3. Repeat 1-2  $k$  times
4. Aggregate predictions of each tree → majority voting

### With and Without Replacement

- With → Chosen feature is put back → can be chosen multiple times
- Without → Chosen feature not put back → feature can't be drawn again

- Random Forest → Advantage → Don't have to worry about choosing good hyperparameters
  - **However** → Doesn't offer same level of interpretability as decision tree
- Don't need to **prune** Random Forest



Only parameter to care about → number of trees,  $k$  → larger number of trees = better performance of random forest classifier → at cost of increased computational cost

- Some hyperparameters can be optimised (less common in practice) → Size,  $n$  + number of features,  $d$ 
  - $n$  controls bias-variance tradeoff

## Bootstrap sample size

- Increasing bootstrap sample → increased diversity among individual trees
- Shrinking bootstrap sample → increase randomness → reduce overfitting → **But**, lower overall performance

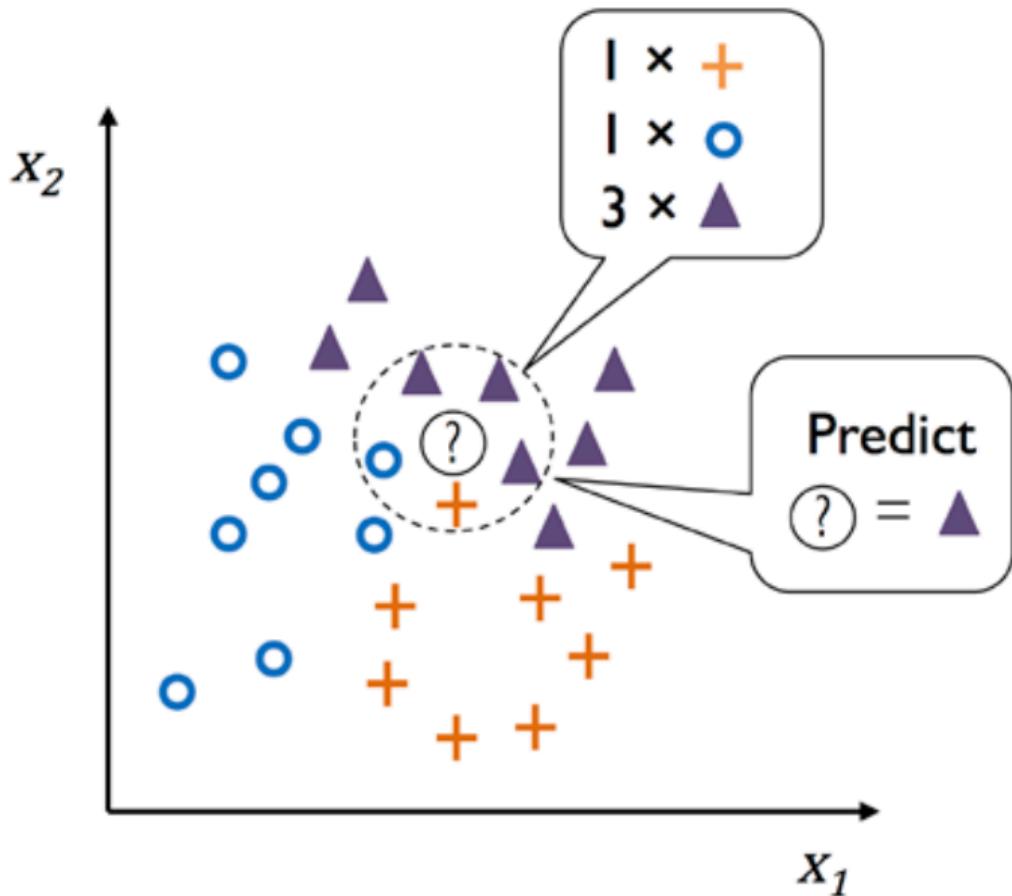
# K-Nearest Neighbors

- K-Nearest Neighbor (KNN) classifier → 'lazy' → not because of its simplicity
  - KNN doesn't learn a discriminative function, memorises training dataset instead

## Parametric vs. Non-parametric models

- Parametric model → **Estimate parameters from training data** → learn function to classify new data points (without need of original training set)
  - **Examples:** Perceptron, Logistic Regression, Linear SVM
- Non-Parametric → **Can't be characterised by a fixed set of parameters + number of parameters change with amount of training data**
  - **Examples:** Decision Tree Classifier/Random Forest, Kernel SVM

- KNN → Sub-category of non-parametric models → **Instance based learning**
  - Memorise datasets
  - Lazy learning in special case → zero cost during learning process



### KNN Algorithm in 3 Steps:

1. Choose number of  $k$  + distance metric
2. Find K-Nearest Neighbor of data that we want to classify
3. Assign new class by **majority vote**

- Finds  $k$  examples in training dataset that are **closest** to the point we want to classify
- Class label of data point → determined by majority vote

### Advantages + Disadvantages of Memory-based Approach

- **Advantage:** New training data → classifier adapts asap.
- **Disadvantage:** Computational complexity grows linearly with number of examples in training dataset → limited storage capabilities
- Efficient data structures for memory based approach = **K-D Tree, Ball Tree**



Most of the times, we'll work with small-medium sized datasets

- Memory based approach → **Good choice for real-world problems**

- **In case of ties** → Scikit-Learn → Prefer neighbor with closer distance to data record to be classified
  - If it's similar lengths → What comes first in order will be chosen first
- Important to find good balance for ***k***
  - Euclidean Distance measure real-value examples
    - **Have to standardise data** so that each feature contributes equally to distance
- **Minkowski Distance** → Generalisation of Euclidean + Manhattan Distance

$$d(x^{(i)}, x^{(j)}) = \sqrt[p]{\sum k |x_k^{(i)} - x_k^{(j)}|^p}$$

- If  $p = 1 \Rightarrow$  Euclidean
- If  $p = 2 \Rightarrow$  Manhattan

## Curse of Dimensionality

- KNN is very susceptible to overfitting → **Curse of Dimensionality**
- Feature space becomes increasingly sparse for increasing number of dimensions of a fixed dataset
- Even your closest neighbors are far away
- **Use feature selection + Dimensionality reduction** → Avoid curse of dimensionality