

Week 5 - ML

Start Date	@1 December 2025
Weeks	Week 5

Data Compression With Dimensionality Reduction - Chapter 5

- Alternative approach to feature selection for dimensionality reduction → Feature extraction
- Data compression important in machine learning → helps store + analyse increasing amount of data that are produced and collected

Unsupervised dimensionality reduction via component analysis

- Feature extraction → transform or project data onto new feature space
- Context of dimensionality reduction → Feature extraction → Data compression with goal to maintain most of the relevant information
 - Not only used to improve storage space or computational efficiency → used to improve predictive performance by reducing Curse of Dimensionality


Principal Component Analysis (PCA)

- PCA → Unsupervised linear transformation technique → used in many different fields → mainly feature extraction + dimensionality reduction
- Some popular application of PCA → Exploratory data analysis, de-noising signals in stock market trading, analysis of genome data + genome expression levels
- PCA → Helps us identify patterns based on correlation between features
 - Looks at high-dimensional data, then asks: "In which directions do the points spread out the most?" → Directions where data varies the most

- Then:

- Builds new axis (new features) along “maximum spread” directions
- Keeps only some of those axes (**often first few**) → end up with fewer dimensions than started with
- Makes sure new axes are at right angles to each other (orthogonal) → new feature captures a different, non-overlapping aspect of variation

SHORT VERSION: PCA rotates + compresses data into fewer, uncorrelated dimensions that capture as much of original variation as possible



Using PCA to find the directions of maximum variance in a dataset

- $x_1, x_2 \rightarrow$ Original feature axes \rightarrow PC1, PC2 \rightarrow Principal components
- PCA for dimensionality reduction $\rightarrow d \cdot x$ - dimensional transformation matrix, W

- Allows us to map vector of feature examples, x to new k -dimensional feature subspace → Fewer dimensions than original d -dimensional feature space

$$x = [x_1, x_2, \dots, x_d], x \in R^d$$

Transformed by transformation matrix, $W \in R^{d \times k}$

$$xW = z$$

Output vector:

$$x = [z_1, z_2, \dots, z_k] = z \in R^k$$

- Result of transforming → first principal component will have largest variance
 - PC1** captures the most information/variation in data → if you look at one axis only → the spread of points is as large as possible
- After **PC1**, PCA looks for second direction (**PC2**) → but, **PC2** must be right angle (orthogonal) to PC1 → so PC1 and PC2 don't overlap.
- This can keep going → **PC3** → has to be orthogonal to PC1 + PC2



PCA directions → sensitive to scaling → Standardise first

Approach Summarised:

- Standardise d -dimensional dataset
- Construct covariance matrix
- Decompose covariance matrix into eigenvectors + eigenvalues
- Sort eigenvalues by decreasing order → to rank corresponding eigenvectors
- Sort eigenvectors → corresponds with k largest eigenvalue → k is dimensionality of new features subspace $K \leq d$
- Construct projection matrix, W from "top" k eigenvectors

- Transform d -dimensional input dataset, \mathbf{X} , using projection matrix, $\mathbf{W} \rightarrow$ obtain new k -dimensional feature subspace

Eigendecomposition

- Factorisation of square matrix into eigenvalues + eigenvectors → core of PCA procedure
- Covariance matrix - $A = A^T$
- Decomposed symmetric vector → Eigenvalues are real numbers + eigenvectors are orthogonal to each other
- Decomposed covariance matrix → Eigenvectors associated with highest eigenvalue corresponds to direction of maximum variance

Four Steps of PCA:

- Standardise the data
- Constructing the covariance
- Obtaining the eigenvalues + eigenvectors of covariance matrix
- Sorting the eigenvalues by decreasing order to rank the eigenvectors

Step 2 of summarised approach

- $d \times d$ -dimensional covariance matrix, d = number of dimensions in the dataset
- Covariance between two features, x_j and $x_k \rightarrow$ on population level:

$$\sigma_{jk} = \frac{1}{n-1} \sum_{i=1}^n (x_j^{(i)} - \mu_j)(x_k^{(i)} - \mu_k)$$

- μ_j, μ_k = sample means of feature j, k
 - Sample mean = zero → if dataset is standardised
- Positive covariance between two features = features increase + decrease together
- Negative covariance → features vary in opposite directions
- Covariance matrix of three features:

$$\Sigma = \begin{bmatrix} \sigma_1^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_2^2 & \sigma_{23} \\ \sigma_{31} & \sigma_{32} & \sigma_3^2 \end{bmatrix}$$

- Covariance matrix represents principal component: (direction of maximum variance)
 - Eigenvalues will define their magnitude

Step 3 of summarised approach

- Obtain eigenpairs for covariance matrix
- Eigenvector v :


$$\Sigma = \lambda v$$

- λ = Scalar \rightarrow eigenvalue

Total and explained variance

- Only select subset of eigenvectors (principal component) that contains most of the information (variance)
- Eigenvalues define magnitude \rightarrow sort by decreasing magnitude \rightarrow interested in top k eigenvectors
- Variance explained ratio

$$\frac{\lambda_j}{\sum_{j=1}^d \lambda_j}$$




The proportion of the total variance captured by the principal components

- First principal → accounts to approx. **40%** of variance
- First two principal components → combined almost **60%** of variance

Feature Transformation

PCA Step 5, 6, 7 of summarised approach

- Sort eigenpairs by descending order of eigenvalues
- Construct projection matrix from selected eigenvectors
- Projection matrix to transform data onto lower-dimensional subspace



Data records from the Wine dataset projected onto a 2D feature space via PCA

- Data more spread along PC1 → Consistent with variance ratio plot (from above)

★ PCA is a unsupervised technique that doesn't use any class label information

Assessing feature contribution

- PCA → create principal components → represent linear combinations of features
- **Loading** → How much each original feature contributes to a given principal component
 - Factor loadings computed → Scaling eigenvectors by square root of eigenvalues
 - **Result:** Correlation between original features + principal components

Supervised Data Compression With Linear Discriminant Analysis (LDA)




Linear Discriminant Analysis (LDA) → Linear transformation technique that takes class label information into account

- General concept → Similar to PCA → LDA finds the feature subspace that optimises class separability



Fisher LDA

- LDA sometimes called Fisher LDA → **Ronald A. Fisher (1936)**
 - Later → Generalised for multi-class problems by **C. Radhakrishnan Rao** → For equal class covariance + normally distributed classes.



The concept of LDA for a two-class problem

- LDA for two-class problem
 - Class 1 - Circles
 - Class 2 - Crosses
- Linear discriminant → x-axis (LD 1) → separate two normal distributed classes
- y-axis (LD 2) → variance in dataset → lot of variance in dataset
 - Would fail at good linear discriminant → doesn't capture any of class-discriminatory information
- **Assumption LDA:**
 - Data is normally distributed
 - Also assumed → classes have identical covariance matrices
 - Training examples are statistically independant.
- If assumptions are slightly violated → LDA for dimensionality reduction can still work reasonably well

LDA Workings:

1. Standardise d -dimensional dataset (d = number of feature)
2. Each class → Compute d -dimensional mean vector → LDA takes labelled information into account
3. Construct between-class scatter matrix S_B + within-class scatter matrix S_W
4. Compute eigenvectors + corresponding eigenvalue of matrix $S_W^{-1}S_B$
5. Sort eigenvalues by decreasing order → rank corresponding eigenvectors
6. Choose k eigenvectors → k largest eigenvalue → construct $d \times k$ -dimensional transformation matrix, W → eigenvector are column of matrix
7. Project examples onto new feature subspace using transformation matrix, W



LDA is similar to PCA → Decomposing matrices into eigenvalues + eigenvectors → For new lower-dimension feature space

Computing Scatter Matrices

- Always standardise the features
- Calculation of new vectors → construct within-class scatter + between-class scatter matrix

$$m_i = \frac{1}{n_i} \sum_{x \in D_i} x_m$$

- m_i = Mean vector
- μ_m = Mean feature value

$$m_i = \begin{bmatrix} \mu_{i\text{alchohol}} \\ \mu_{i\text{malicacid}} \\ \vdots \\ \mu_{i\text{proline}} \end{bmatrix}^T \quad i \in \{1, 2, 3\}$$

- Mean vectors → Compute within-class scatter matrix, S_W :

$$S_W = \sum_{i=1}^c S_i$$

- Calculated by summing up individual scatter matrices:

$$S_i = \sum_{x \in D_i} (x - m_i)(x - m_i)^T$$

- Assumption: Class labels in training datasets are uniformly distributed.
 - Violated → Check with class label distribution

Scale individual scatter matrices, S_i → Before we sum them up.

- Divide scatter matrices by number of class-examples, n_i
 - Scatter matrix → Same as computing covariance matrix Σ_i

- Normalised version of scatter matrix

$$\Sigma_i = \frac{1}{n_i} S_i = \frac{1}{n_i} \sum_{x \in D_i} (x - m_i)(x - m_i)^T$$

- Next steps → compute between-class scatter matrix S_B :

$$S_B = \sum_{i=1}^c n_i (m_i - m)(m_i - m)^T$$


Linear Discriminants For New Feature Subspace

- Instead of eigendecomposition → Solve generalised eigenvalue problem of matrix $S_W^i S_B$
- LDA → Number of linear discriminants at most $c - 1$, c = Number of class labels → S_B = Sum of c matrices with rank one or less



Collinearity → Perfect case (All aligned example points fall on straight line)

- Covariance matrix = rank one → eigenvector with nonzero eigenvalue



The top two discriminants capture 100 percent of the useful information


- Measure how much class-discriminatory information captured by linear discriminants (eigenvectors)
- First two linear discriminants = capture 100% of useful information

```
Matrix W:
[[ -0.1481 -0.4092]
 [ 0.0908 -0.1577]
 [-0.0168 -0.3537]
 [ 0.1484  0.3223]
 [-0.0163 -0.0817]
 [ 0.1913  0.0842]
 [-0.7338  0.2823]
 [-0.075   -0.0102]
 [ 0.0018  0.0907]
 [ 0.294   -0.2152]
 [-0.0328  0.2747]
 [-0.3547 -0.0124]
 [-0.3915 -0.5958]]
```

- Two discriminative eigenvector columns stacked → create transformation matrix W


Projecting Examples Onto New Feature Subspace

- Transformation matrix $W \rightarrow$ Transform training dataset \rightarrow Multiply matrices
 $\rightarrow X' = XW$




Wine classes perfectly separable after projecting the data onto the first two discriminants

- Wine classes \rightarrow perfectly linearly separable in new feature subspace



The logistic regression model misclassifies one of the classes

- Lower-dimension dataset → logistic regression
 - Only one example is misclassified



The logistic regression model works perfectly on the test data

- Lower regularisation strength → shifts decision boundaries and all test data is separated perfectly

Nonlinear Dimensionality Reduction via Visualisation

- **T-distributed Neighbor Embedding (t-SNE)** → Frequently used in literature
→ Visualise high-dimensional datasets into two or three dimensions
- t-SNE → Plot image of handwritten images in 2-dimensional space

Why consider Nonlinear dimensionality Reduction?


- Development + application of nonlinear dimensionality reduction technique
→ **Manifold Learning** → Scikit-Learn library covers other complex ones too




Manifold Learning → Lower dimensional topological space embedded in high-dimensional space

- Algorithm has to capture complicated structure of data → to project it onto lower-dimensional space → Relationship between data points is preserved

Different views of a 3-dimensional Swiss roll:



Swiss roll projected onto a 2-dimensional feature space with ...



Three-dimensional Swiss roll projected into a lower, two-dimensional space

- Very powerful but also very hard to use**
 - Difficult → Working with high-dimensional datasets that we cannot visualise + structure is not obvious
 - Could try projecting dataset to 2 or 3 dimensions
 - Not enough for capturing complicated relationships
 - Also for it's hard or impossible to assess quality



Hence people rely on PCA + LDA for dimensional reduction

Visualising Data via t-Distributed Stochastic Neighbor Embedding

- In a nutshell → t-SNE modelling data points based on pair-wise distances in high-dimensional (original) feature space
 - Then → t-SNE learns to embed data points into a lower-dimensional space → pair-wise distances in original space is preserved
- t-SNE → Intended for visualisation purpose → Requires whole dataset for projection
 - Projects points directly → Cannot apply t-SNE to new data points

Uniform Manifold Approximation And Projection

- Another popular visualisation technique → UMAP
- Faster than t-SNE + produces similarly good results as t-SNE + used to project new data