

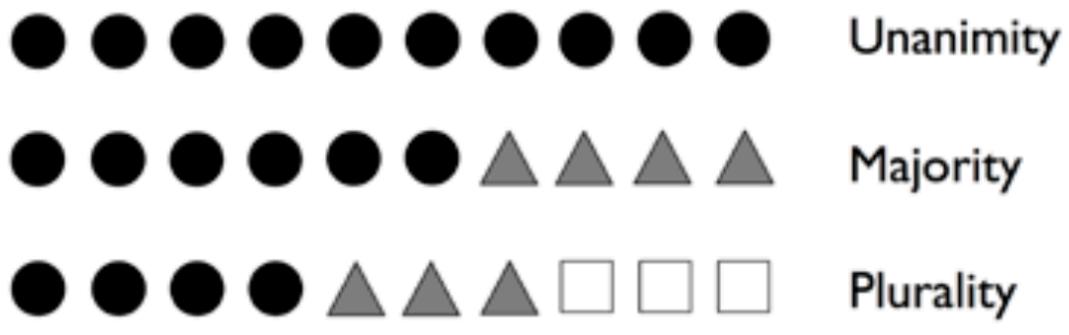
# Week 7 - ML

Start Date	@15 December 2025
Weeks	Week 7

## Combining Different Models For Ensemble Learning

### Learning With Ensemble

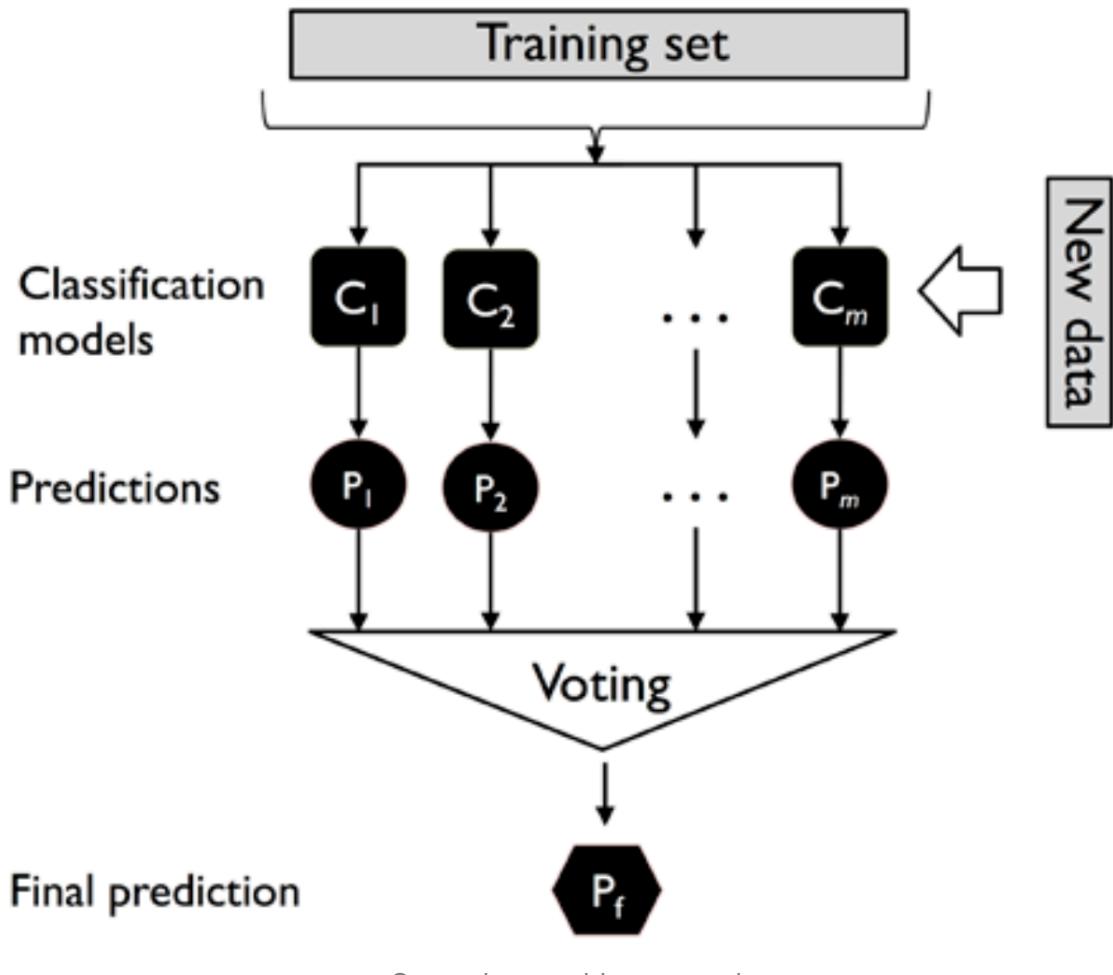
- Ensemble method → Combine classifiers to meta-classifier → Better generalisation performance
- Majority Voting Principle → Class label with more than 50% of votes → refers to binary class settings only
  - Generalise to multi-class settings → Plurality Voting (UK → 'Absolute' + 'Relative' Voting) → Class with most votes



Different voting concepts

- Training dataset → Training  $m$  different classifiers
- Depending on technique
  - Use different classification algorithm
  - Decision Trees, SVM, Logistic Regression

- Can also use same base classification algorithm, which fits different subsets of training datasets



- Predict class label with Majority or Plurality voting, then **combine class labels classifier**,  $C_j$  + **selected class label**,  $\hat{y}$

$$\hat{y} = mode \{C_1(x), C_2(x), \dots, C_m(x)\}$$

- This shows the most frequent event. (mode from stats)
- Binary Classification + Majority Vote:

$$C(x) = sign \left[ \sum_{j=1}^m C_j(x) \right] = \begin{cases} 1, & \text{if } \sum_{j=1}^m C_j(x) \geq 0, \\ -1, & \text{otherwise.} \end{cases}$$

- Ensemble method works better than individual classifiers → **Proven with combinatorics**

- Assumptions for Binary Classification task:
  - $n$ -base classifiers → equal error rate,  $\varepsilon$
  - Classifiers independent + error rate not correlated

$$P(y_g \geq k) = \sum_k^n \binom{n}{k} \varepsilon^k (1 - \varepsilon)^{n-k} = \varepsilon_{\text{ensemble}}$$

- $\binom{n}{k}$  = Binomial Coefficient
- Error probability of ensemble of base classifiers as probability mass function of binomial distribution
- Example = 11 Classifiers ( $n = 11$ ) + Error rate 0.25 ( $\varepsilon = 0.25$ )

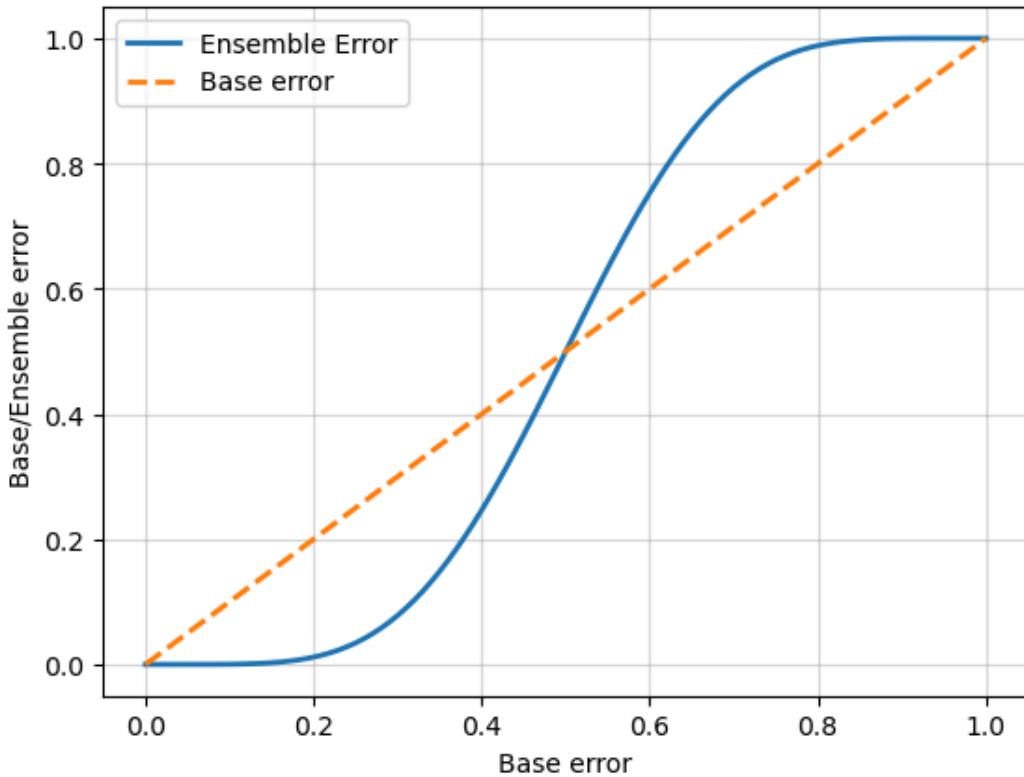
$$P(y \geq k) = \sum_{k=6}^{11} \binom{11}{k} 0.25^k (1 - 0.25)^{11-k} = 0.034$$

- Error rate lower than of each individual classifier

## Binomial Coefficient

- Number of ways we choose  $k$  from set of size  $n$  → "n choose k"

$$\frac{n!}{(n - k)!k!}$$



Plot of the ensemble error versus the base error

- Error probability of ensemble always better

## Combining Classifiers via Majority Vote

### Single Majority Vote Classifier

$$\hat{y} = \arg_i \max \sum_{j=1}^m w_j \chi_A(C_j(x) = i)$$

- $w_j$  = Weight associated with base classifier,  $C_j$
- $\hat{y}$  = Predicted class label of the ensemble
- $A$  = Set of unique class labels
- $\chi_A$  = Characteristic or indicator function
- For equal weights  $\rightarrow$  simplified equation:

$$\hat{y} = \text{mode} \{C_1(x), C_2(x), \dots, C_m(x)\}$$

- Example:

- $C_1(x) = 0, C_2(x) = 0, C_3(x) = 1$
- $\hat{y} = \text{mode } \{0, 0, 1\} = 0$
- Weights =  $C_1, C_2 = 0.2, C_3 = 0.6$

$$\hat{y} = \arg \max [0.2 \times i_0 + 0.2 \times i_0, 0.6 \times i_1] = 1$$

- $3 \times 0.2 = 0.6 \Rightarrow C_3$  has three times the weight than  $C_1$   $C_2 \rightarrow \hat{y} = \text{mode } \{0, 0, 1, 1, 1\} = 1$
- Modified version of Majority Vote for predicting class labels (using `predict_proba` method from Logistic Regression):

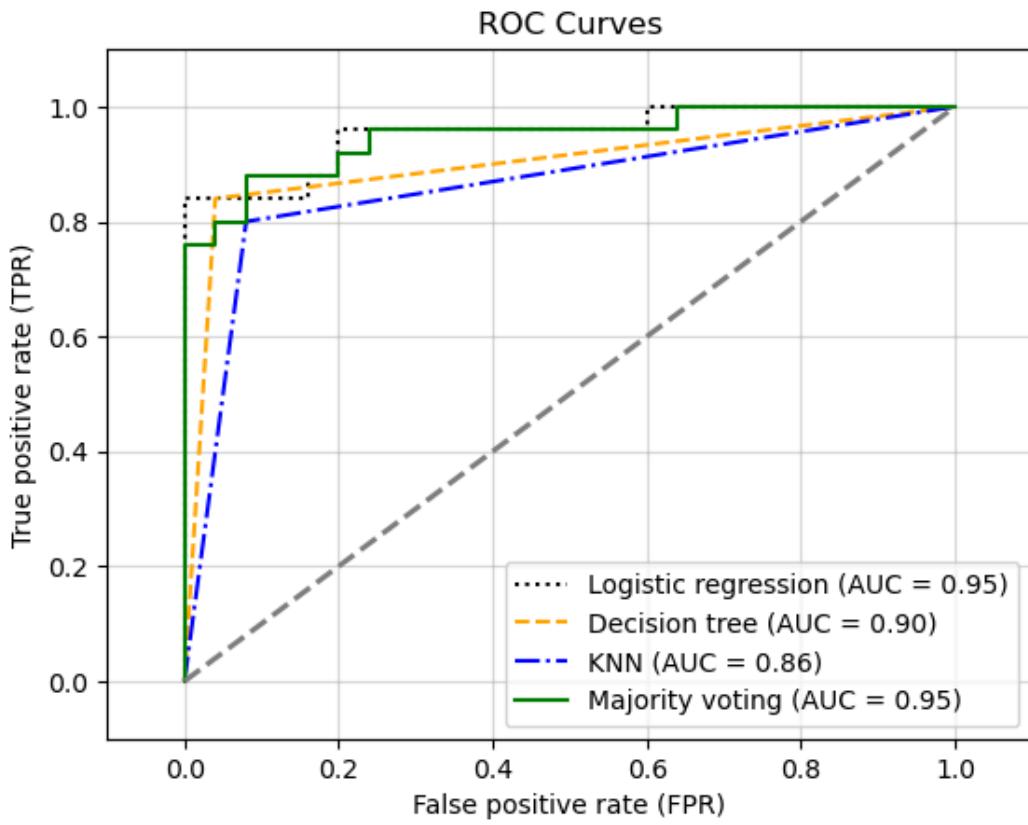
$$\hat{y} = \arg \max \sum_{j=1}^m w_j P_{ij}$$

- $P_{ij}$  = Predicted probability of  $j$ th class label  $i$

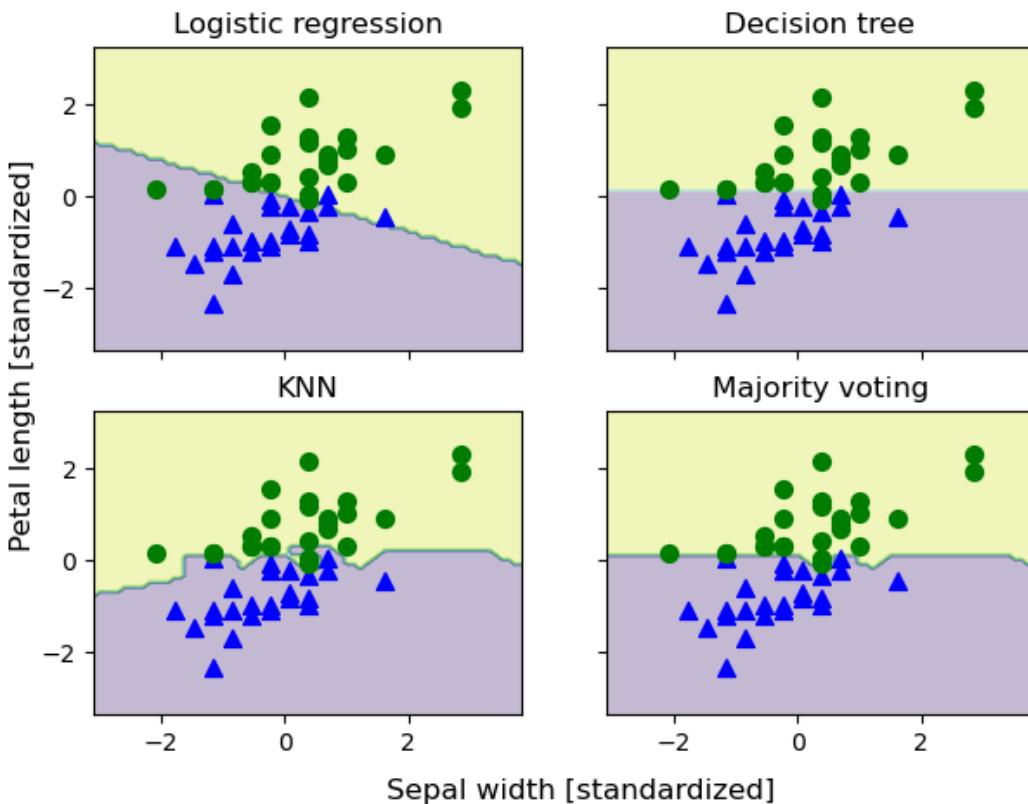
- Binary Classification problem + same weights

$$C_1(x) \rightarrow [0.9, 0.1], \quad C_2(x) \rightarrow [0.8, 0.2], \quad C_3(x) \rightarrow [0.4, 0.6]$$

$$\begin{aligned} p(\hat{y}_0 | x) &= 0.2 \times 0.9 + 0.2 \times 0.8 + 0.6 \times 0.4 = 0.58 \\ p(\hat{y}_1 | x) &= 0.2 \times 0.1 + 0.2 \times 0.2 + 0.6 \times 0.6 = 0.42 \\ \hat{y} &= \arg \max_i [p(\hat{y}_0 | x), p(\hat{y}_1 | x)] = 0 \end{aligned}$$



The ROC curve for the different classifiers



The decision boundaries for the different classifiers

# Stacking (Ensemble Methods)

- **David H. Wolpert** → "Stacked Generalization" 1992

**Input:** Training data  $D = \{\chi_i, y_i\}_{i=1}^n (x_i \in R^n, y_i \in \gamma)$

**Output:** Ensemble classifier  $H$

## Step 1: Learn first-level classifiers

for  $t \leftarrow 1$  to  $T$  do

    Learn a base classifier  $h_t$  based on  $D$

end for

## Step 2: Construct new data sets from $D$

for  $i \leftarrow 1$  to  $n$  do

    Construct a new data set that contains  $\{x'_i, y_i\}$  where  $x' = \{h_1(x_i), h_2(x_i), \dots, h_T(x_i)\}$

- $x'$  = Modified feature vector
- $\{h_1(x_i), h_2(x_i), \dots, h_T(x_i)\}$  = predicted class labels via scikit-learn

end for

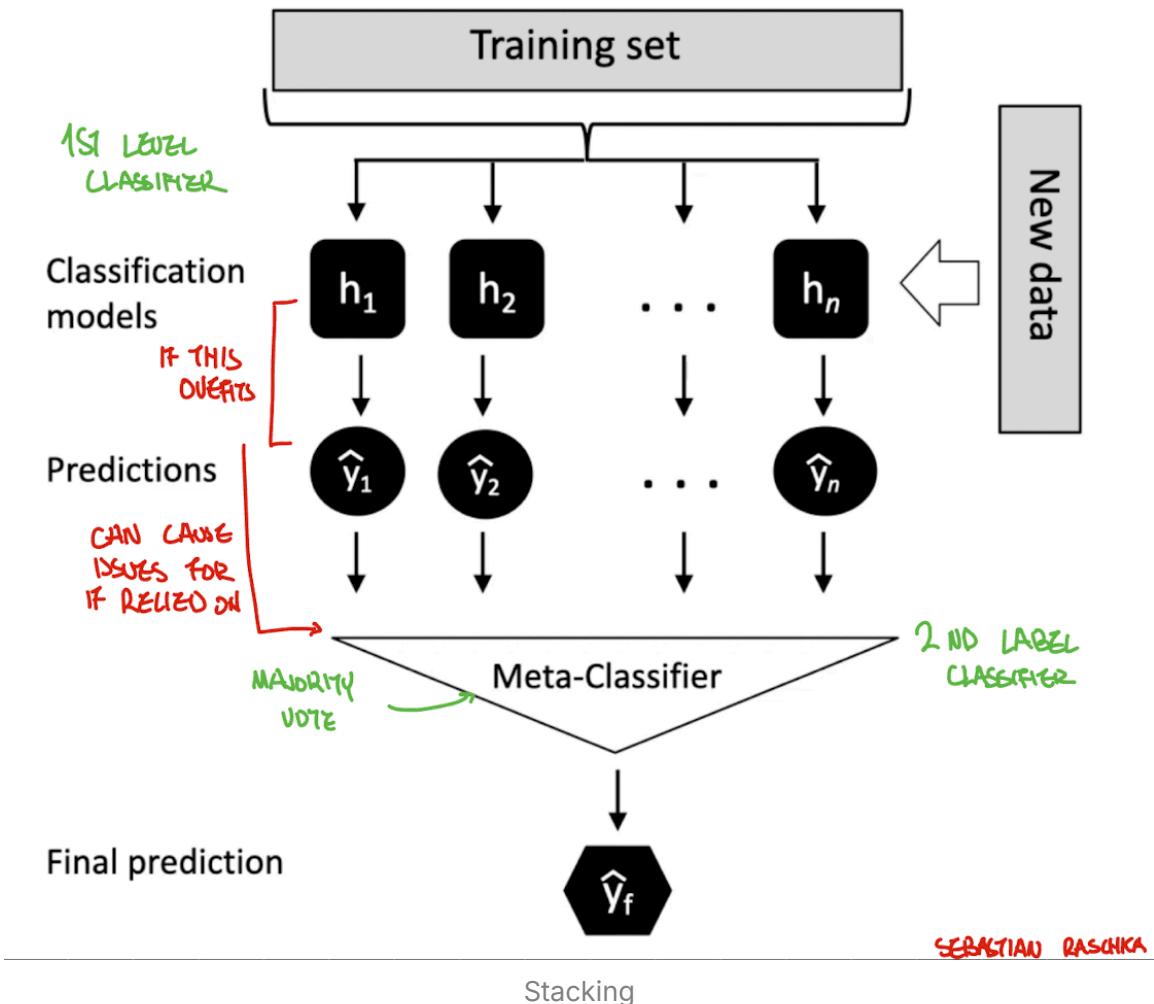


The first two steps are the same as Majority Voting

## Step 3: Learn a second-level classifier → Learn new classifier based on predictions

Learn a new classifier  $h'$  based on newly constructed data set.

return  $H(x) = h'(h_1(x), h_2(x), \dots, h_T(x))$



- **Problem with stacking??**
  - Prone to overfitting...
- You can improve stacking with cross-validation
- Stacking → K-fold classification for 2nd label classifier
  - Use k predictions for 2nd label classifier

## Modified Stacking Algorithm

**Input + Output are the same**

### Step 1: Cross validation approach preparation for 2nd-level classifier

Randomly split  $D$  into  $k$  equal-sized subsets  $\rightarrow D = \{D_1, D_2, \dots, D_k\}$   
 for  $k \leftarrow 1$  to  $K$  do

#### Step 1.1: Learn 1st-level classifiers

```

for  $t \leftarrow 1$   $T$  do
    Learn classifier  $h_{KT}$  from  $D/D_k$ 
end for

```

### Step 1.2: Construct training set for 2nd-level classifier

```

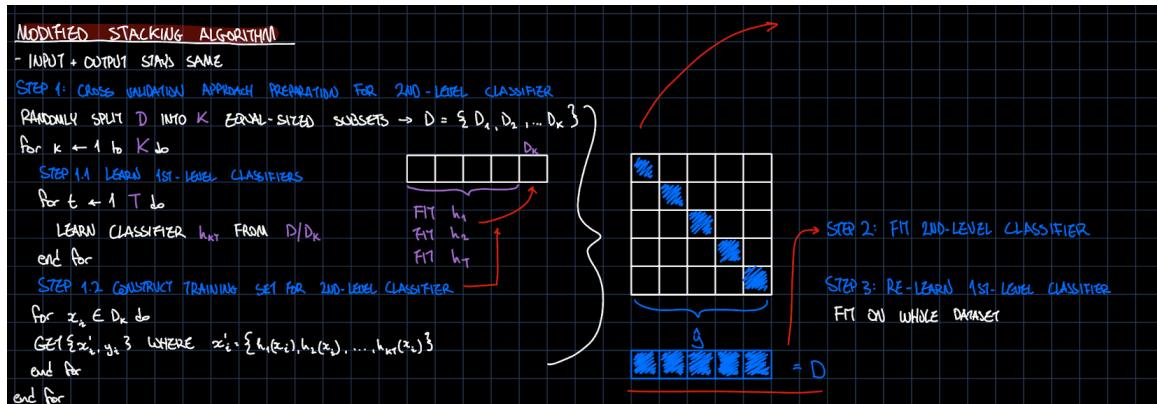
for  $x_i \in D_k$  do
    Get  $\{x'_i, y_i\}$  where  $x'_i = \{h_1(x_i), h_2(x_i), \dots, h_{KT}(x_i)\}$ 
end for

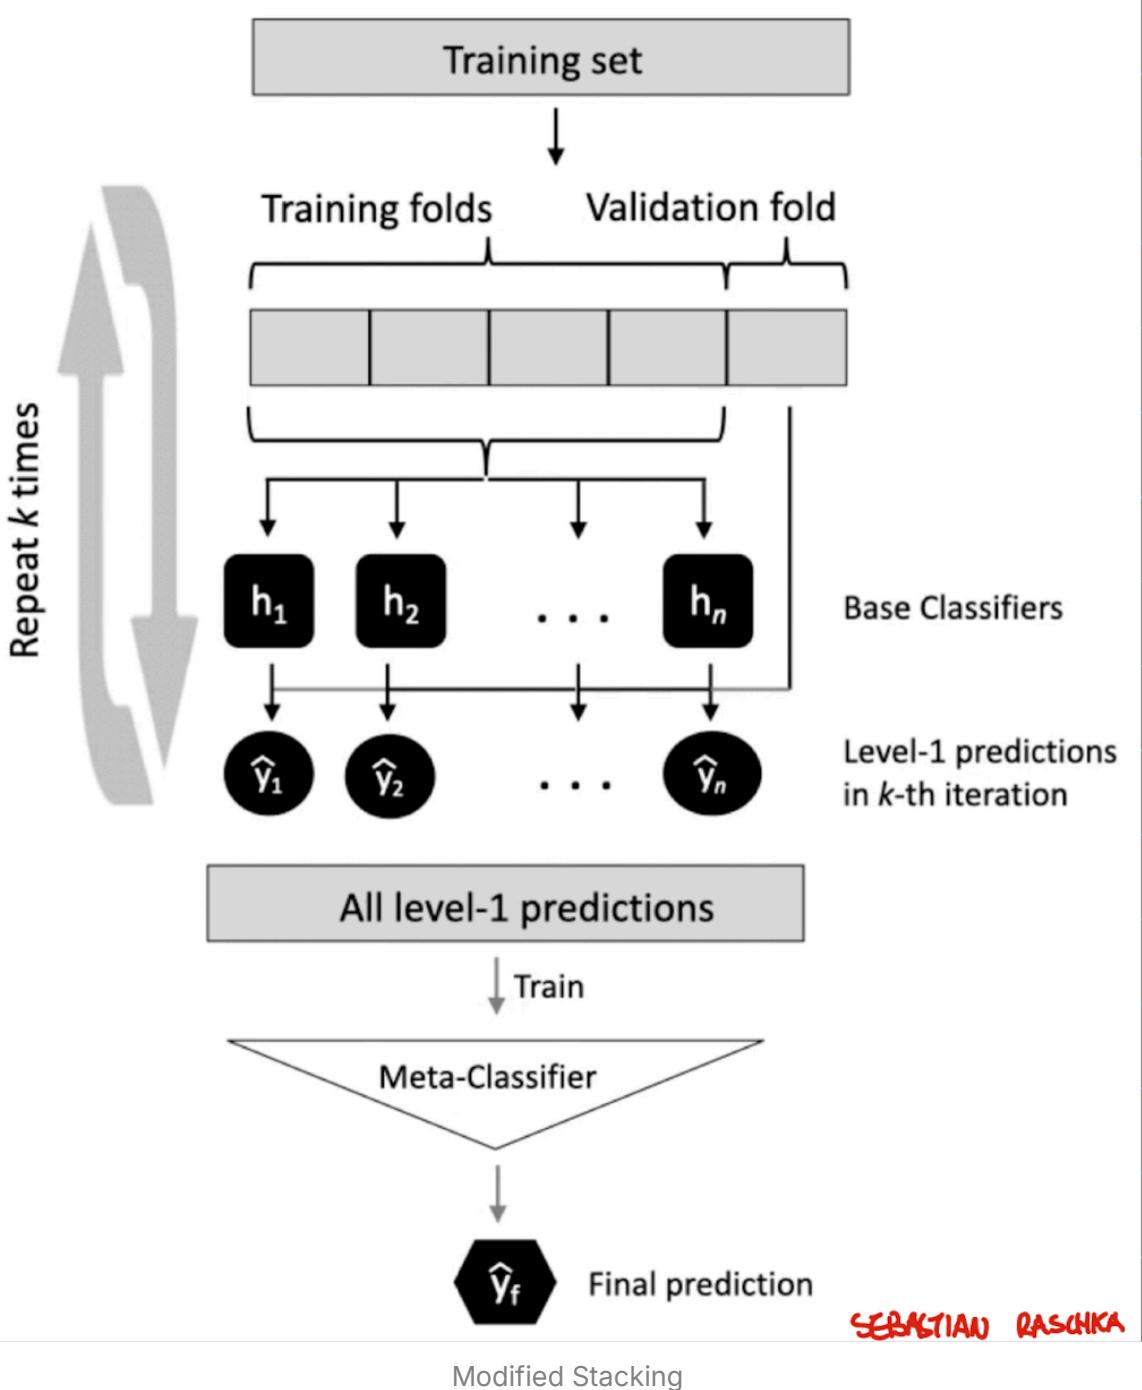
```

end for

### Step 2: Fit 2nd-level classifier

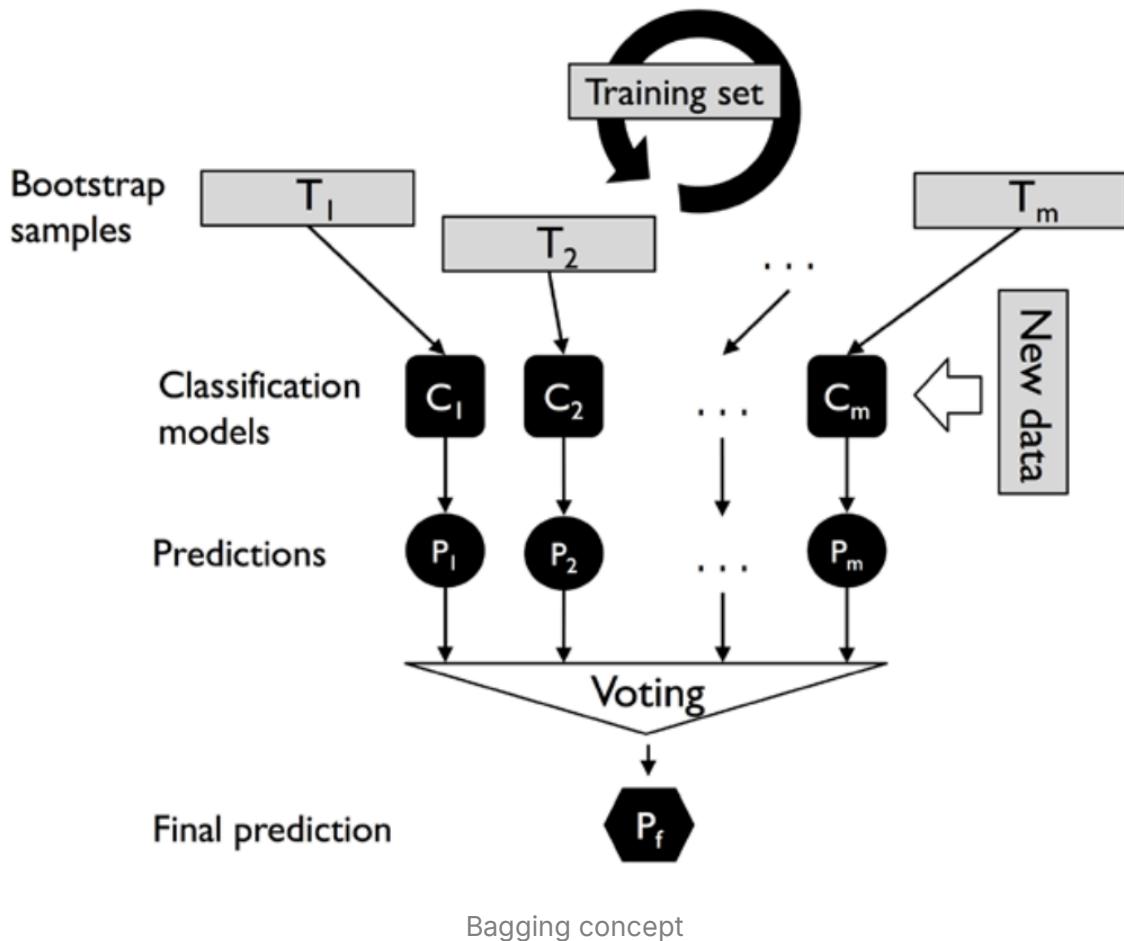
### Step 3: Re-learn 1st-level classifier



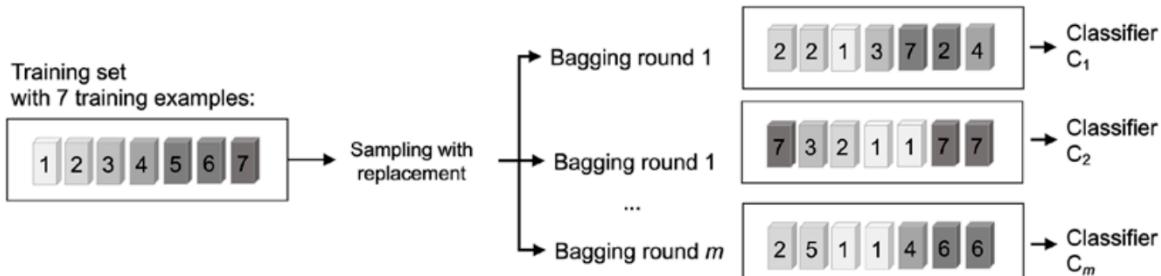


## Bagging - Ensemble Classifiers From Bootstrap Samples

- Closely related to `MajorityVoteClassifier`
- Instead of same training dataset to fit individual classifiers → Draw bootstrap samples (random samples with replacement) from initial training dataset
  - Called **Bootstrap Aggregating**



Bagging concept



Example of bagging

- Random samples obtained via Bagging → Bagging round
- Each subset contains some duplicates + some original examples don't appear
  - **Sampling with replacement**
- Individual classifiers fit to bootstrap samples and then predictions combined with Majority Vote
- Bagging related to Random Forest Classifier

- Random Forest → Special example → Random feature subsets also used when fitting individual decision trees
- Complex classification tasks + dataset's high dimensionality can lead to overfitting in single decision trees → In these scenarios bagging comes in handy



**Bagging is an effective approach to reducing variance of model**

- However, it's ineffective in reducing model bias, if models too simple to capture trends in data
  - This is why we want to perform bagging on ensemble of classifiers with low bias → e.g. Unpruned Decision Trees

## Adaptive Boosting - Leveraging weak learners

- Adaptive boosting (AdaBoost)
  - Robert E. Shapire (1990) → 'The Strength of Weak Learnability'
- Boosting - Ensemble consists of very simple base classifiers → Referred to as 'weak learners' → Slight edge over random guessing
  - Key Concept → Focus on training examples that are hard to classify
    - Weak learners subsequently learn from misclassified examples to improve performance of ensemble

### How Adaptive Boosting works:



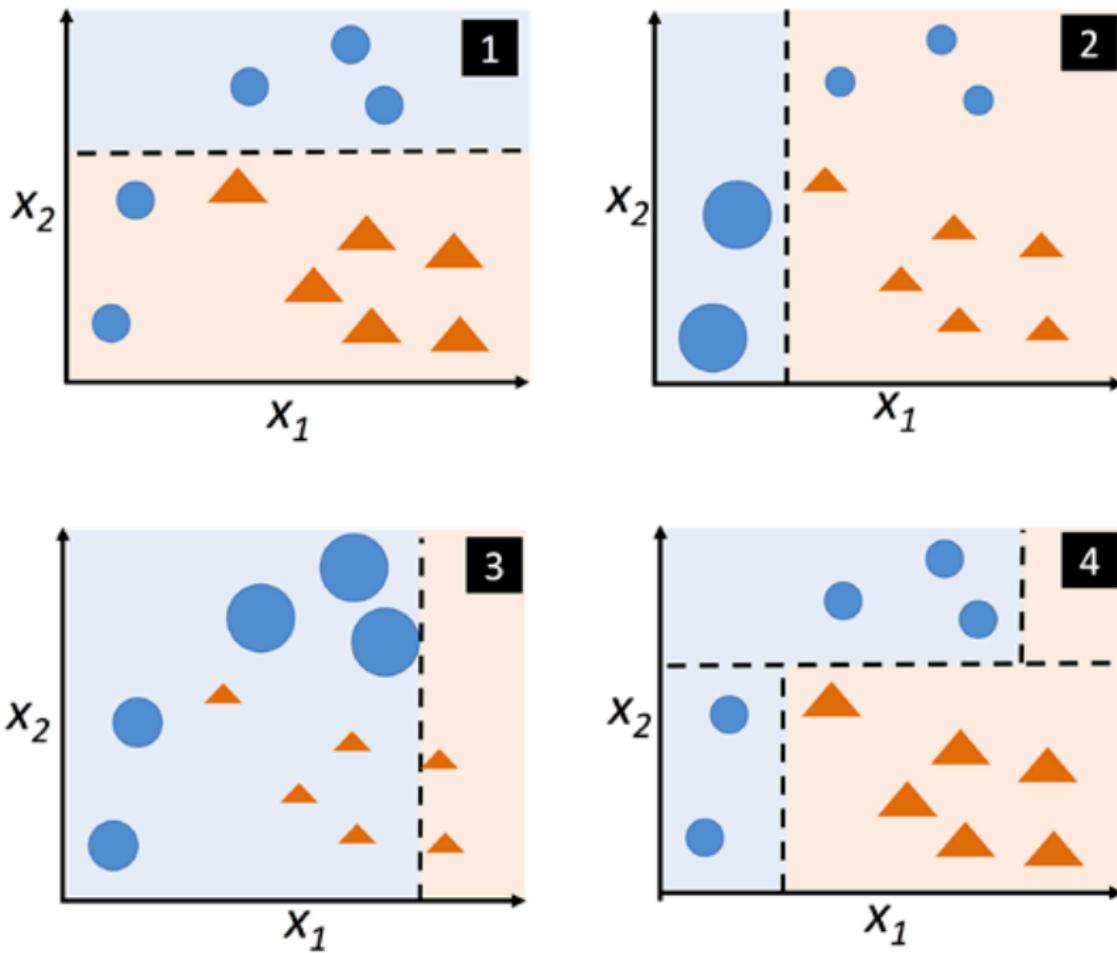
**Initial formulation - Algorithm uses random subsets of training examples from training dataset without replacement**

### Four key steps:

1. Draw random subset (samples) from training examples,  $d_1$  without replacement from training dataset,  $D$  for weaker learner  $C_1$
2. Second random training subset,  $d_2$  without replacement and 50% of examples from previously misclassified to train weaker learner  $C_2$

3. Find training examples,  $d_2$  in training dataset,  $D$  which  $C_1 + C_2$  disagree with, to train  $C_3$
4. Combine  $C_1, C_2, C_3$  via Majority Voting

- **Leo Breiman** → Boosting can lead to decrease in bias + variance - compared to bagging
  - In practice, AdaBoost known for high variance → tend to overfit training data
- AdaBoost - uses the complete training dataset to train weak learners → training examples reweighed after each iteration to learn from mistakes



The concept of AdaBoost to improve weak learners

### AdaBoost steps:

1. Set weight vector,  $w$ , to uniform weights,  $\sum_i w_i = 1$

2. For  $j$  in  $m$  boosting rounds:

- Train weighted weak learner,  $C_j = \text{train}(X, y, w)$
  - Predict class labels,  $\hat{y} = \text{predict}(C_j, X)$
  - Compute weighted error rate,  $\epsilon = w \cdot (\hat{y} \neq y)$
  - Compute coefficient,  $\alpha_j = 0.5 \log\left(\frac{1-\epsilon}{\epsilon}\right)$
  - Update weights,  $w := w \times \exp(-\alpha_j \times \hat{y} \times y)$
  - Normalise weights to sum to 1,  $w := w / \sum_i w_i$
3. Compute final prediction,  $\hat{y} = \left( \sum_{j=1}^m (\alpha_j \times \text{predict}(C_j, X)) > 0 \right)$



**Considered bad practice to select a model based on the repeated usage of the test dataset. → Generalisation performance may be overoptimistic**



- Ensemble learning increases computational complexity compared to individual classifiers
  - In practice, ask yourself, is it worth the modest improvement in predictive performance for increased computational cost

## Boosting vs. Stacking

	Boosting	Stacking
<b>Training Order</b>	Sequential training	Parallel training
<b>Core Goal</b>	Reduce bias by repeatedly correcting mistakes	Exploit diversity - learn how to mix different strong models
<b>New Model Fitting</b>	Emphasises misclassified/large error samples	Each base model sees same data, meta-model trained on predictions
<b>Combination Method (final predictions)</b>	Weighted sum / large-error samples	Meta-model output

<b>Parallelisation</b>	Limited (Iteration dependant)	Base learners easily parallelisable
<b>Sensitivity</b>	Can overfit/noise-sensitive if not regularised	Can overfit if meta-model is too flexible → relies on good CV setup
<b>Typical Use</b>	Gradient Boosting Trees, XG Boost, AdaBoost	Competition/production 'blender' → Combine Trees, Linear Models, SVMs, Neural Nets, etc...

## Gradient Boosting - Training Ensemble Based on Loss Gradients



**Gradient boosting is important → Forms basis of popular ml algo like XGBoost → Well known for kaggle competitions**

- Another variant of boosting concept → Training weaker learners
- **Gradient boost** fits decisions trees in an iterative fashion using prediction errors
  - Deeper than decision tree stumps and typically a maximum depth of 3 to 6 (max 8-64 leaf nodes)
  - Does not use prediction errors for assigning sample weights, instead it goes directly to form target variable fro fitting the next tree
  - Uses same global learning rate for each tree

## Outline of Gradient Boost Algorithm



**Gradient boost is a general-purpose supervised learning method**

- Here we'll look gradient boost in a classification model (binary classification example)
- Builds series of trees, where each tree is fit on error (difference between label and predicted value)

- Each round, tree ensemble improves, which nudges the tree more in the right direction with small updates
- Updates based on loss gradient, that's how it got the gradient 'boosting' name
- **Step-by-step (general algorithm):**

1. Initialise model to return constant prediction value. Decision tree root node (decision tree with single leaf node)

$$F_0(x) = \arg_{\hat{y}} \min \sum_{i=1}^n L(y_i, \hat{y})$$

- $\hat{y}$  = value returned by decision tree
- $L$  = loss function
- $n$  =  $n$  training examples

2. For each tree  $m = 1, \dots, M$ ,  $M$  is user-specified total number of trees

- a. Compute difference between predicted value  $F(x_i) = \hat{y}_i$  and class label  $y_i$ , sometimes called **Pseudo-Response** or **Pseudo-Residual**

**Formally → Write Pseudo-residual as negative of loss function with respect to predicted values:**

$$r_{im} = \left[ \frac{\partial(y_i, F(x_i))}{\partial F(x_i)} \right]_{F(x)=F_{m-1}(x)}$$

- $F(x)$  = prediction of previous tree  $F_{m-1}(x)$
- **If it's the first round then the constant value from single node tree instead of prediction of the previous tree**
- c. Fit a tree to pesudo-residual  $r_{im}$ ,  $R_{jm}$  to denote  $j = 1, \dots, J_m$  lead nodes of resulting tree in iteration  $m$
- d. For each lead node  $R_{jm}$  → compute output value:

$$\gamma_{jm} = \arg_{\gamma} \min \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma)$$

- $\gamma_{jm}$  = computed by

minimising  
loss function



$R_{jm}$  can contain more than one training example, hence summation

- e. Update model by adding output value  $\gamma_m$  to previous tree:

$$F_m(x) = F_{m-1}(x) + \eta \gamma_m$$

Instead of adding full predicted values of current tree  $\gamma_m$  to previous tree  $F_{m-1}$ , scale  $\gamma_m$  by learning rate  $\eta$  (small step → typically between 0.01 - 1)

## Gradient Boost for Classification

- Single training example, logistic loss:

$$L_i = -y_i \log p_i + (1 - y_i) \log(1 - p_i)$$

- $\log(\text{odds})$ :

$$\hat{y} = \log(\text{odds}) = \log\left(\frac{p}{1-p}\right)$$

- Use  $\log(\text{odds})$  to rewrite equation:

$$L_i = \log(1 + e^{\hat{y}_i}) - y_i \hat{y}_i$$

- Partial derivative of loss function with respect to  $\log(\text{odds})$ ,  $\hat{y}$ :

$$\frac{\partial L_i}{\partial \hat{y}_i} = \frac{e^{\hat{y}_i}}{1 + e^{\hat{y}_i}} - y_i = p_i - y_i$$

After this is complete, the next step is to now add the gradient boost steps:

1. Create root node, which will minimise logistic loss. Loss minimised if root node return  $\log(\text{odds})$ ,  $\hat{y}$
2. For each tree  $m = 1, \dots, M$ ,  $M$  is user-specified total number of trees

- a. Convert  $\log(\text{odds})$  into probability using logistic function (we used this in logistic regression):

$$p = \frac{1}{1 + e^{-y}}$$

Then we compute pseudo-residual, which is the **negative partial derivative of loss with respect to  $\log(\text{odds})$**

- This is the difference between class label and the predicted probability
- b. Fit new tree to pseudo-residuals
- c. For each  $R_{jm}$ , compute value  $\gamma_{jm}$ , which minimises loss function

$$\begin{aligned}\gamma_{jm} &= \arg\min \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma) \\ &= \log(1 + e^{\hat{y}_i + \gamma}) - y_i(\hat{y}_i + \gamma)\end{aligned}$$

**This results in:**

$$\gamma_{jm} = \frac{\sum_i y_i - p_i}{\sum_i p_i(i - p_i)}$$

- This is only for  $R_{jm}$ , and not full training set
- d. Update model by adding gamma value from  $2c$  with learning rate  $\eta$ :

$$F_m(x) = F_{m-1}(x) + \eta \gamma_m$$



**After boosting algorithm is complete → predict class labels by thresholding probability values of final model,  $F_m(x) = 0.5$ , like you would do in logistic regression**

- **However, boosting will produce non-linear decision boundaries unlike logistics regression**

## XGBoost

- Gradient boosting is a sequential process, which is slow to train (inefficient time wise)
- The alternative is **XGBoost**
- XGBoost (Extreme Gradient Boost) → Proposed several tricks + approximations → speeds up the training process → used a lot in kaggle competitions